

MATH-454 Parallel and High Performance Computing Lecture 1: Execution in an HPC environment

Pablo Antolin

Slides of N. Richart, E. Lanti, V. Keller's lecture notes

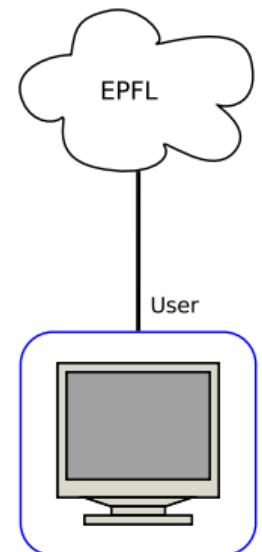
February 27 2025

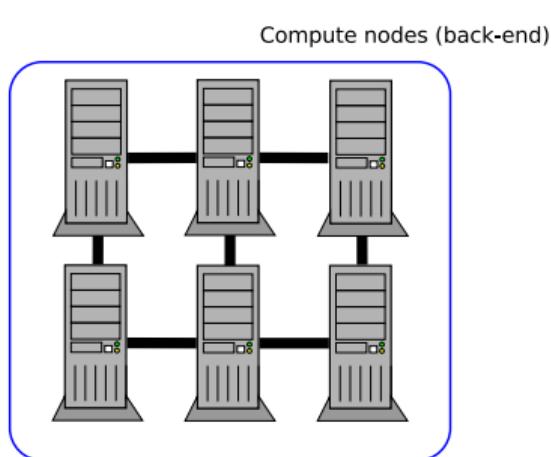
Today exercise session

- What is a cluster.
- How to connect to a cluster (any remote machine).
- How to compile codes and use libraries.
- How to submit job on a cluster.
- How to use GIT

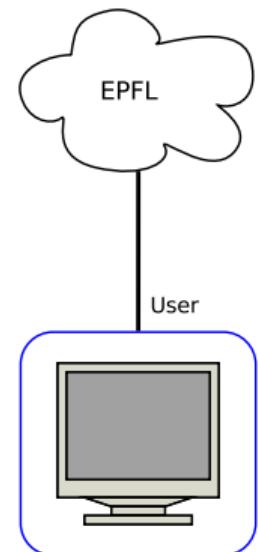
groups.epfl.ch

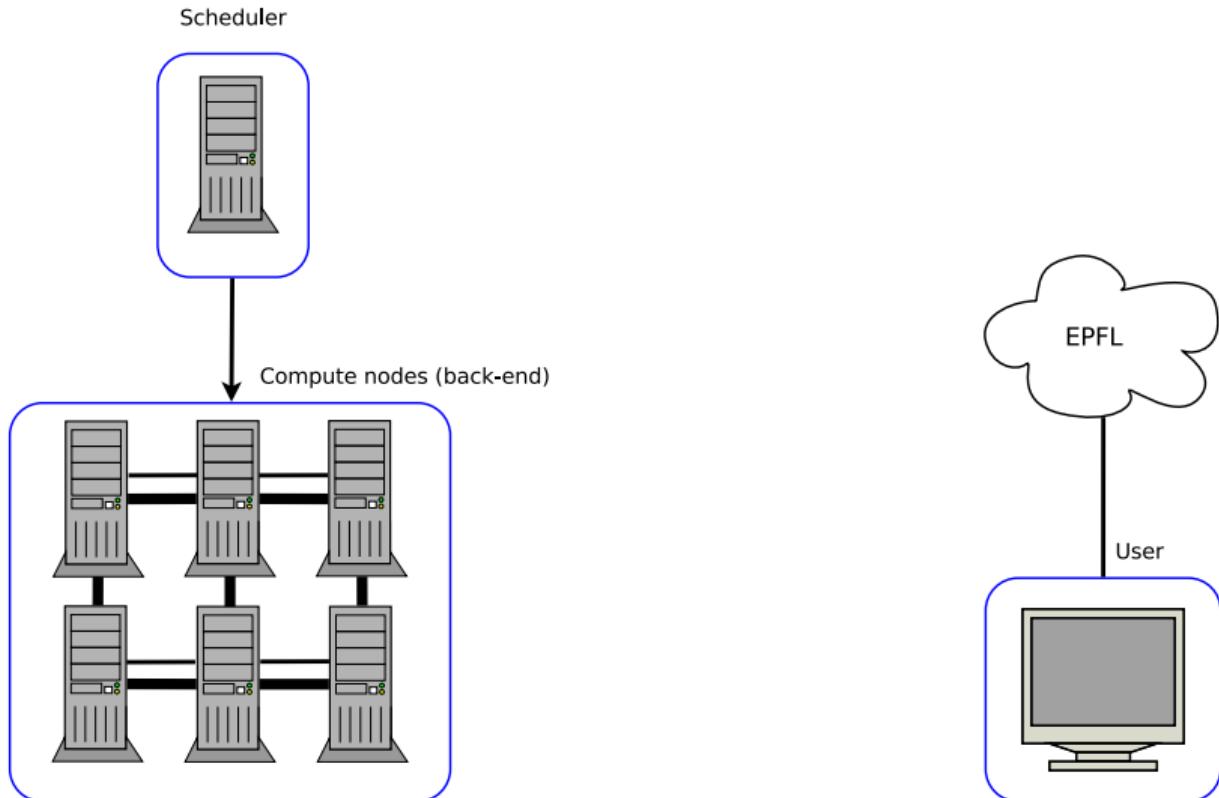
If you are not in the group [hpc-math454](#) let me know

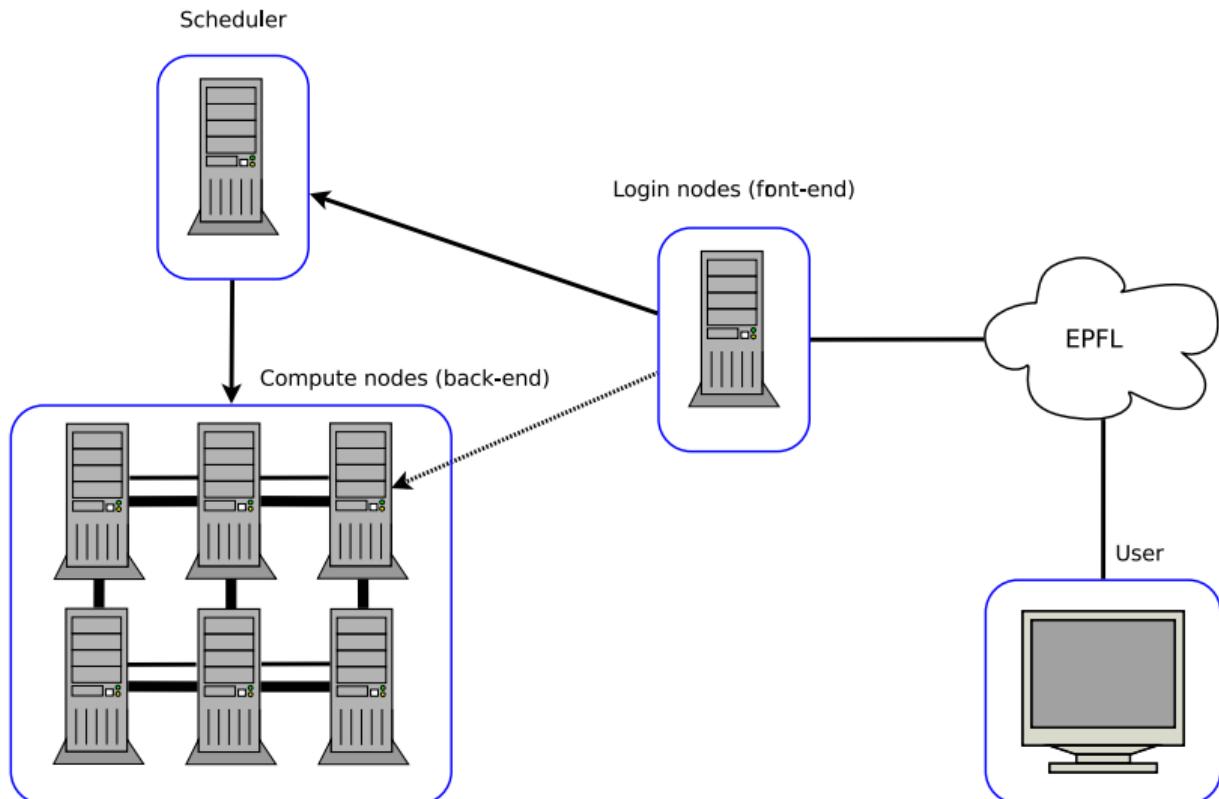


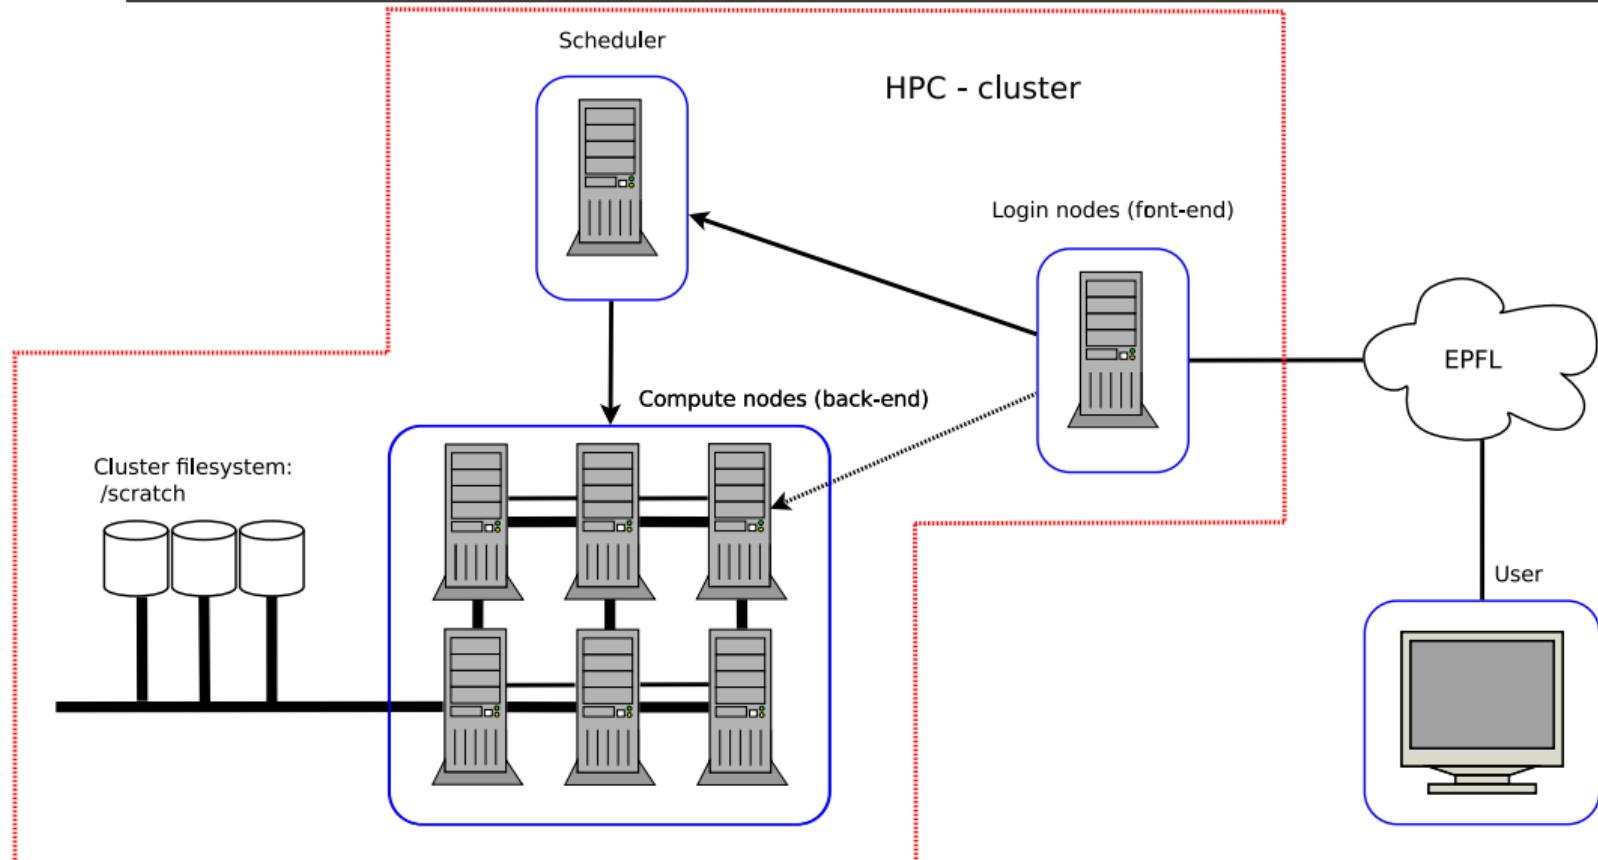


P. Antolin

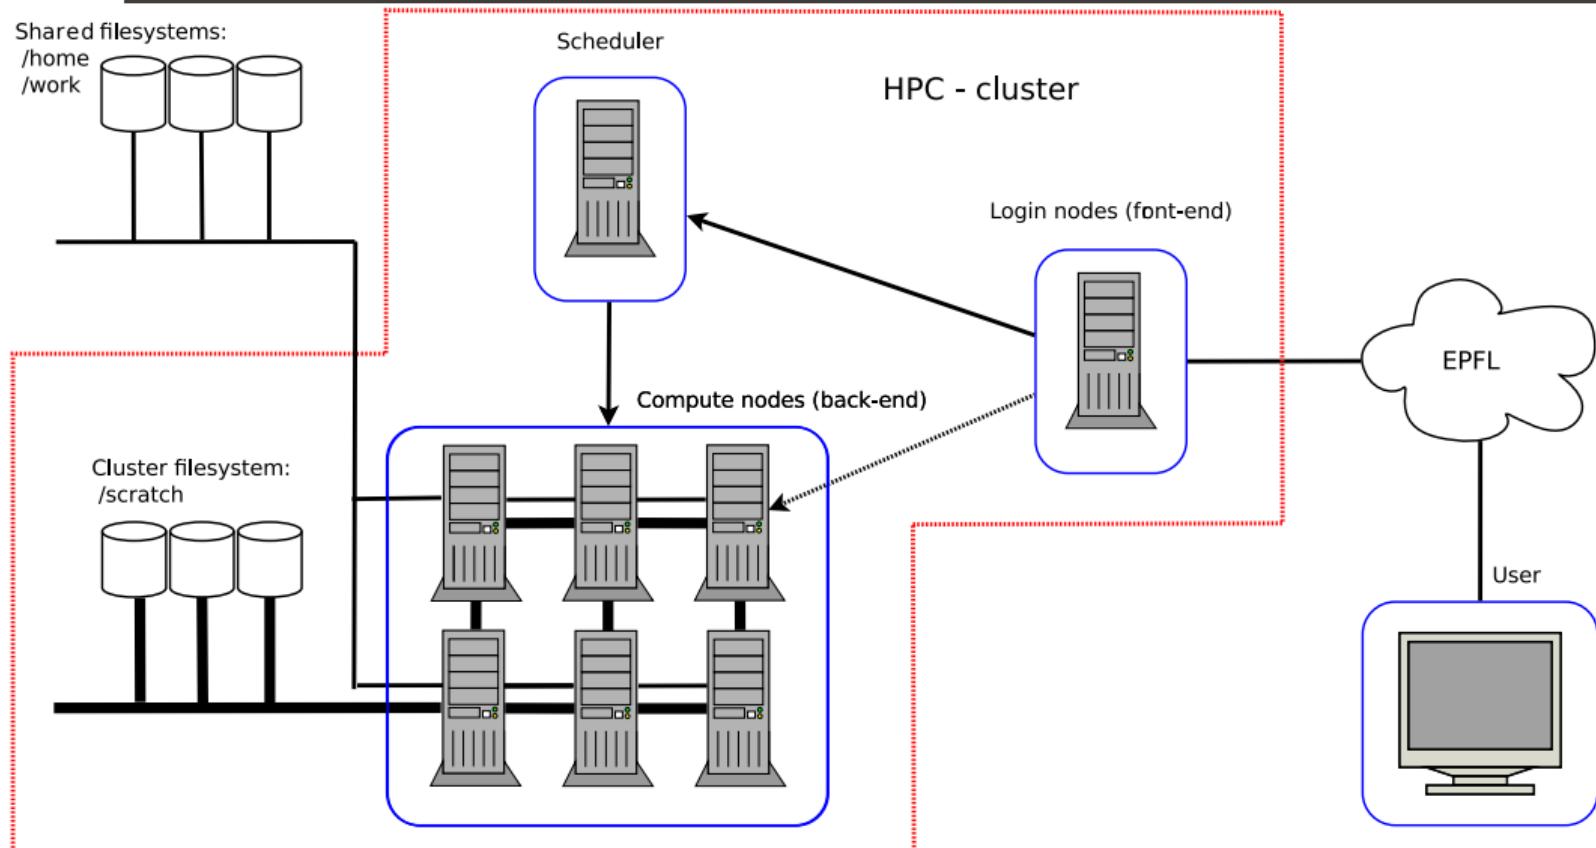








What is a cluster?



Login	Hosts	Cores	RAM	Network
hostname	#	# x GHz	GB	Gbit/s
helvetios.hpc.epfl.ch	287	36x2.3	192	100 (IB)
jed.hpc.epfl.ch	375	72x2.4	512	Ethernet
	42	72x2.4	1024	Ethernet
	2	72x2.4	2048	Ethernet

Login	Hosts	Cores	RAM	Network
hostname	#	# x GHz	GB	Gbit/s
izar.hpc.epfl.ch	35	40x2.1	192	2x100 (IB) + 2x NVidia V100 7TFlops
	35	40x2.1	384	100 (IB) + 4x NVidia V100 7.8TFlops
	2	40x2.1	768	100 (IB) + 4x NVidia V100 7.8TFlops

Login	Hosts	Cores	RAM	Network
hostname	#	# x GHz	GB	Gbit/s
kuma.hpc.epfl.ch	84	32x2.7	384	2x200 (IB) + 4x NVidia H100 94 GB
	20	32x2.7	384	2x200 (IB) + 8x NVidia L40S 48 GB
	2	32x2.7	384	2x200 (IB)

- The simulation data is written on the storage systems. At SCITAS:
 - ▶ **/home**: store source files, input data, small files
 - ▶ **/work**: collaboration space for a group
 - ▶ **/scratch**: temporary huge result files
- Please, note that only **/home** and **/work** have backups!
- **/scratch** data can be erased at any moment!

First step

- Connect to a remote cluster to get a shell
- SSH: Secure SHell

How to use

```
$ ssh -l <username> <hostname>
$ ssh <username>@<hostname>
```

First step

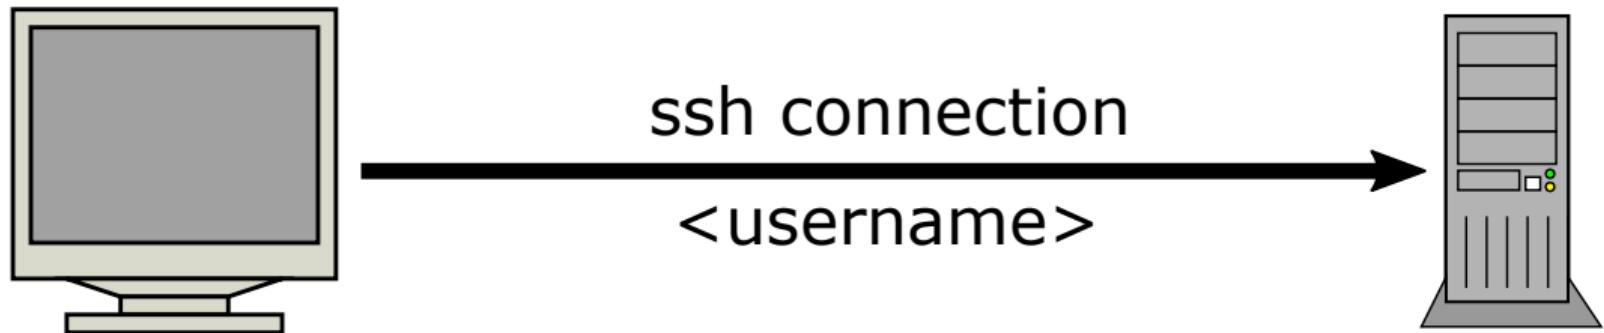
- Connect to a remote cluster to get a shell
- SSH: Secure SHell

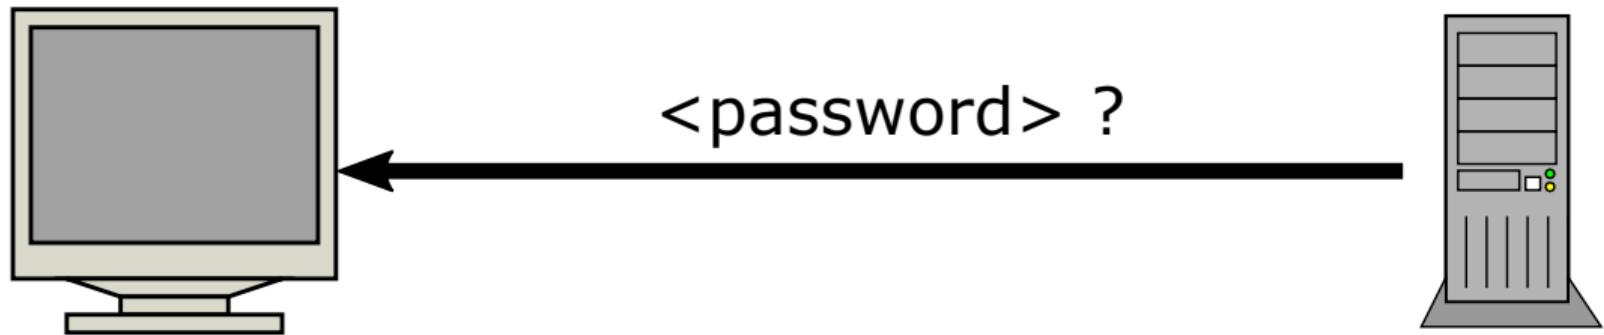
How to use

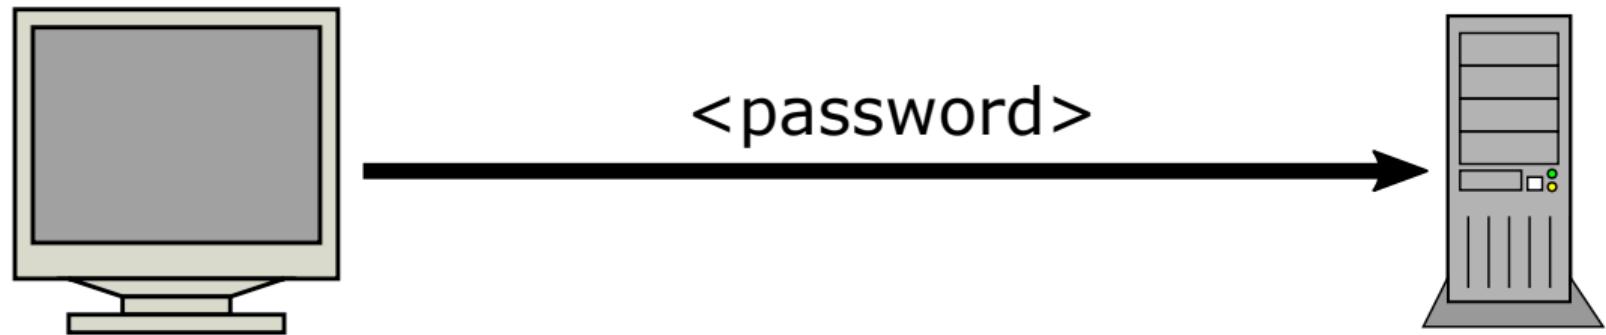
```
$ ssh -l <username> <hostname>  
$ ssh <username>@<hostname>
```

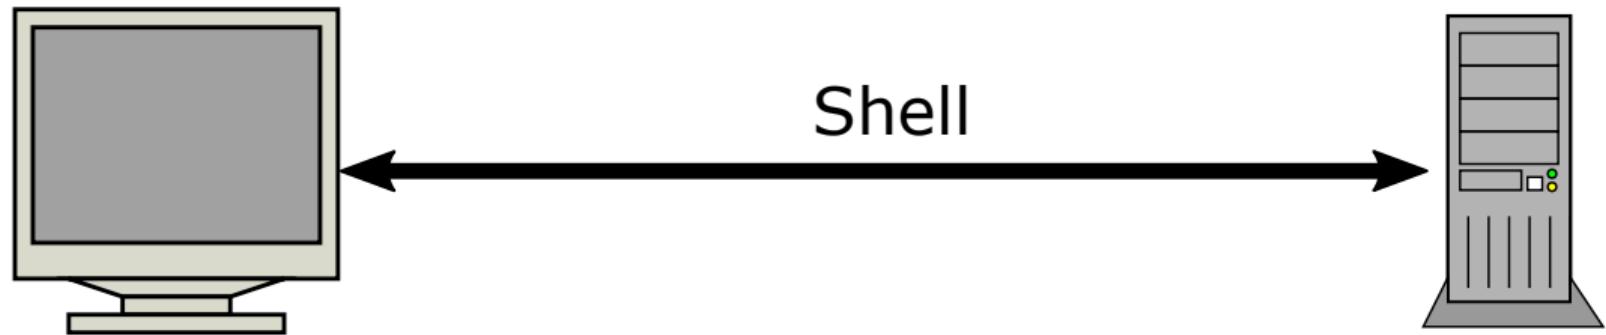
For windows users

Just install git, and use git bash









To connect to the front node of a cluster

```
$ ssh -l jdoe helvetios.hpc.epfl.ch  
$ ssh jdoe@helvetios.hpc.epfl.ch
```

Front nodes

- **helvetios** [*CPU (OpenMP/MPI)*]
 - **izar** [*GPU (CUDA)*]
-
- Connect to helvetios' front node
 - Check the different folders **/home /scratch**

scp works fine for Linux/MacOS, it is often pre-installed
On windows use GIT Bash or pscp.exe from Putty

How to use scp/pscp.exe

Send data to remote machine:

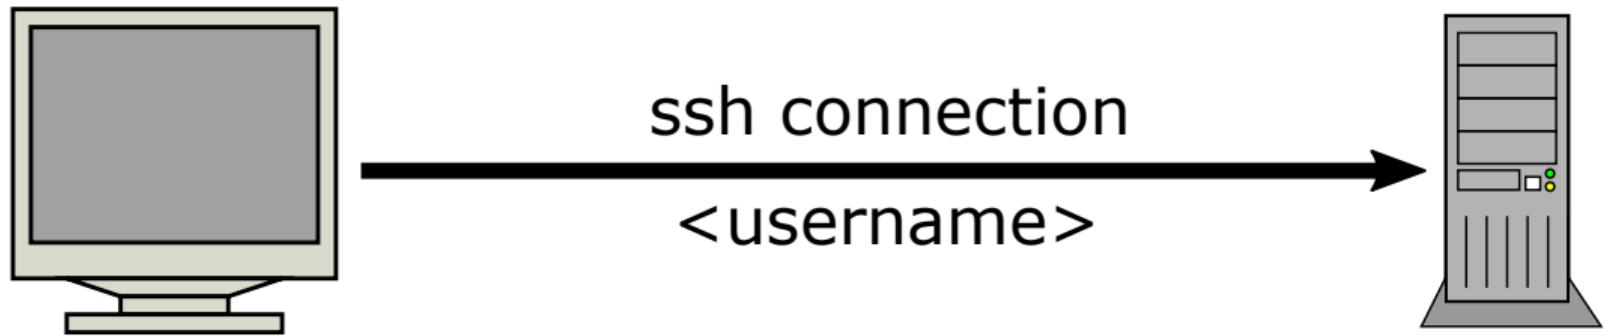
```
$ scp [-r] <local_path> <username>@<remote>:<remote_path>
```

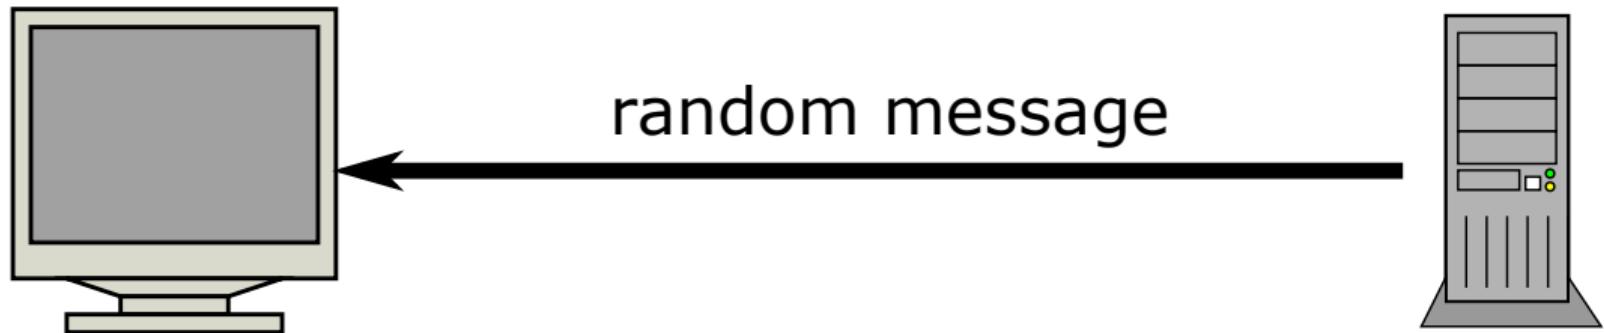
Retrieve data from remote machine:

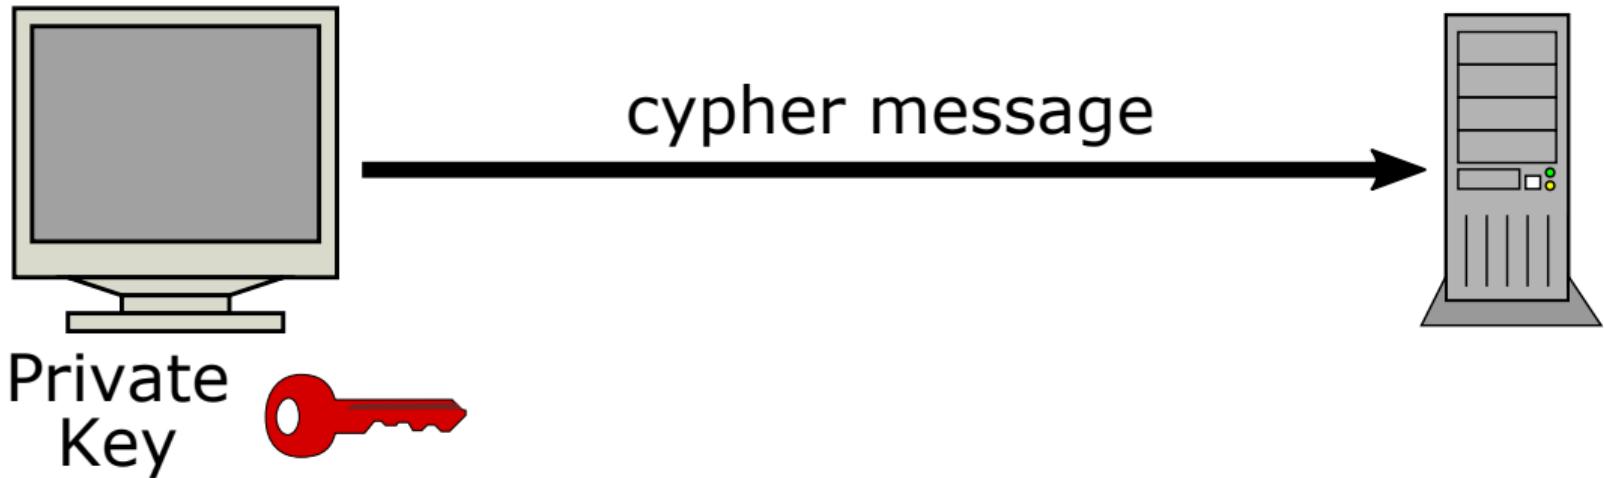
```
$ scp [-r] <username>@<remote>:<remote_path> <local_path>
```

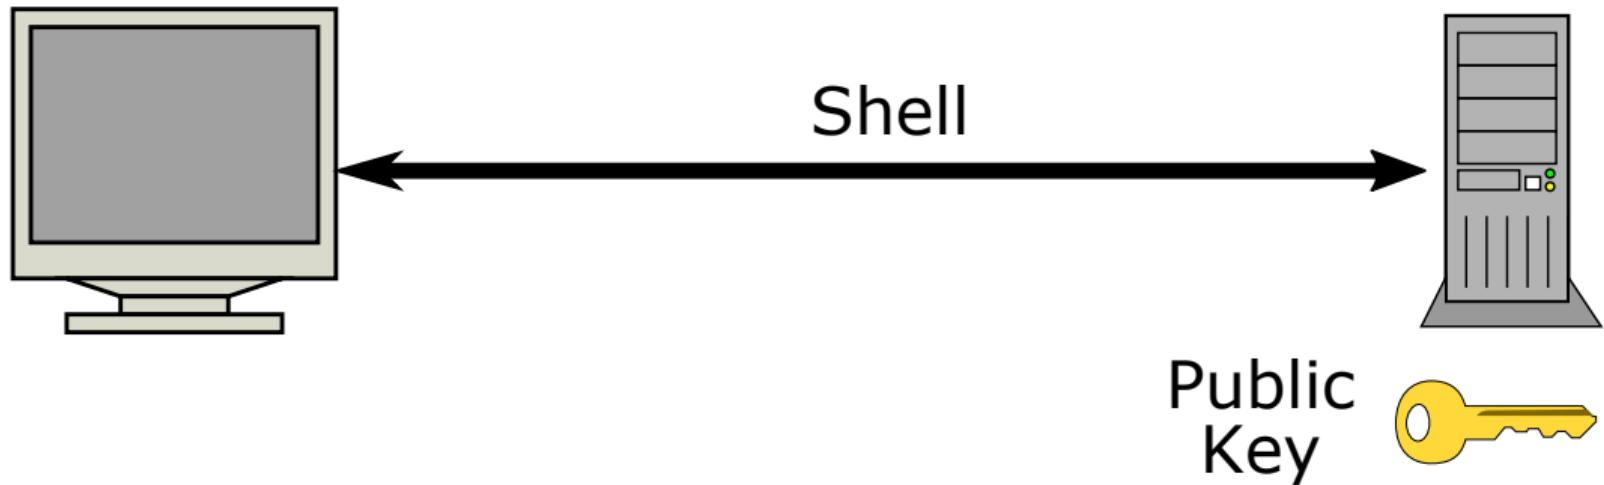
Note: It is always easier to “send to” and “receive from” the clusters since they have a fix ip/name

- Copy a file from your machine to the cluster.
- Modify the file (e.g., using **vim** or **nano**) and retrieve it from the cluster onto your machine









- Generate a pair of public/private key using `ssh-keygen`.
- Copy the generated public key
- Try to connect (it should not ask your password)
- Same ssh key should be already working for `helvetios`, `jed`, and `izar`

Example

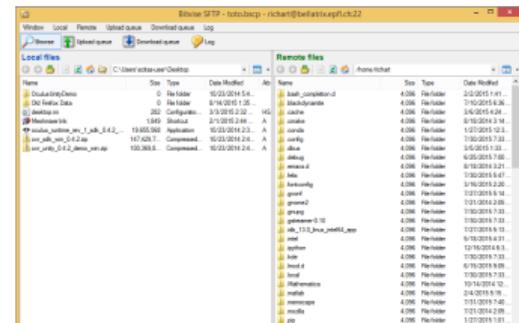
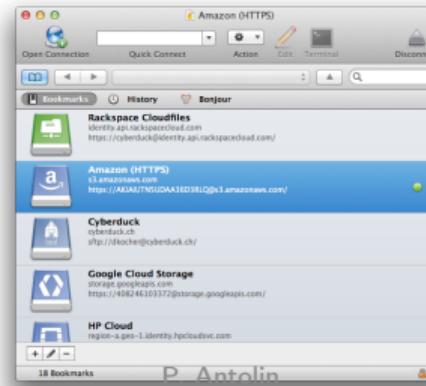
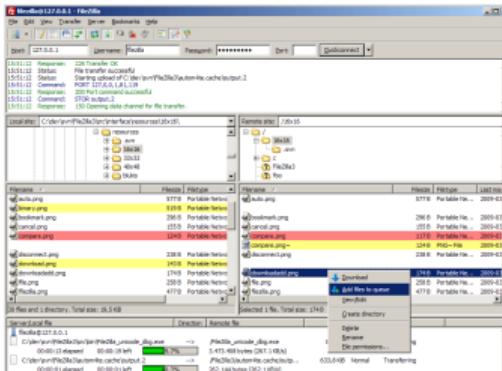
```
$ ssh-keygen -b 4096 -t rsa
[ Follow instructions ]
$ ssh-copy-id jdoe@helvetios.hpc.epfl.ch
```

sftp tones of GUIs, for example: Filezilla <https://filezilla-project.org> ("all" OSes)

Cyberduck <https://cyberduck.io/> (Windows and Mac)

Tunnelier <https://bitvise.com/tunnelier> (Windows)

TUIs also exists like lftp



It is common practice to use VS Code to edit code on remote machines just as if it was your local machine.

Steps:

1. Install the "Remote - SSH" Extension
 - ▶ Open VS Code
 - ▶ Go to Extensions (Ctrl+Shift+X on Windows/Linux, Cmd+Shift+X on Mac)
 - ▶ Search for "Remote - SSH" and install it
2. Open the Command Palette (Ctrl+Shift+P on Windows/Linux, Cmd+Shift+P on Mac)
3. Select: Remote-SSH: Connect to Host...
4. Enter your remote machine details:
 - ▶ Example: ssh username@helvetios.hpc.epfl.ch
5. Authenticate
 - ▶ Enter your password or use an SSH key as shown before
6. Start coding!

What are modules

- A way to dynamically modify the environment
- They contain configurations to use an application/a library

How to use them

- **module avail** list all possible modules
- **module load <module>** load a module
- **module unload <module>** unload a module
- **module purge** unload all the modules
- **module list** list all loaded modules
- **module spider <module>** info about a module

- List all available modules
- Try **g++ --version**
- Load the module **gcc**
- Try again **g++ --version**
- Try **icpc --version**
- Load the **intel** module
- Try again **icpc --version**
- List the currently loaded modules

GCC:

```
$ g++ -o <executable> <sources>
```

Intel:

```
$ icpc -o <executable> <sources>
```

Compilation all-in-one

If the compiler is given source code only it will do all the stages at once

- Load a compiler module
- Compile the code hello.cc

- Generate object files: compile

```
$ g++ -c file1.cc  
$ g++ -c file2.cc
```

Transform code file `file?.cc` in object file `file?.o`

- Generate the executable: linking

```
$ g++ -o hello file1.o file2.o
```

Transforms object file `file?.o` in executable `hello`

- Translation is made by a compiler in 4 steps

Preprocessing Format source code to make it ready for compilation (remove comments, execute preprocessing directives such as `#include`, etc.)

Compiling Translate the source code (C, C++, Fortran, etc) into assembly, a very basic CPU-dependent language

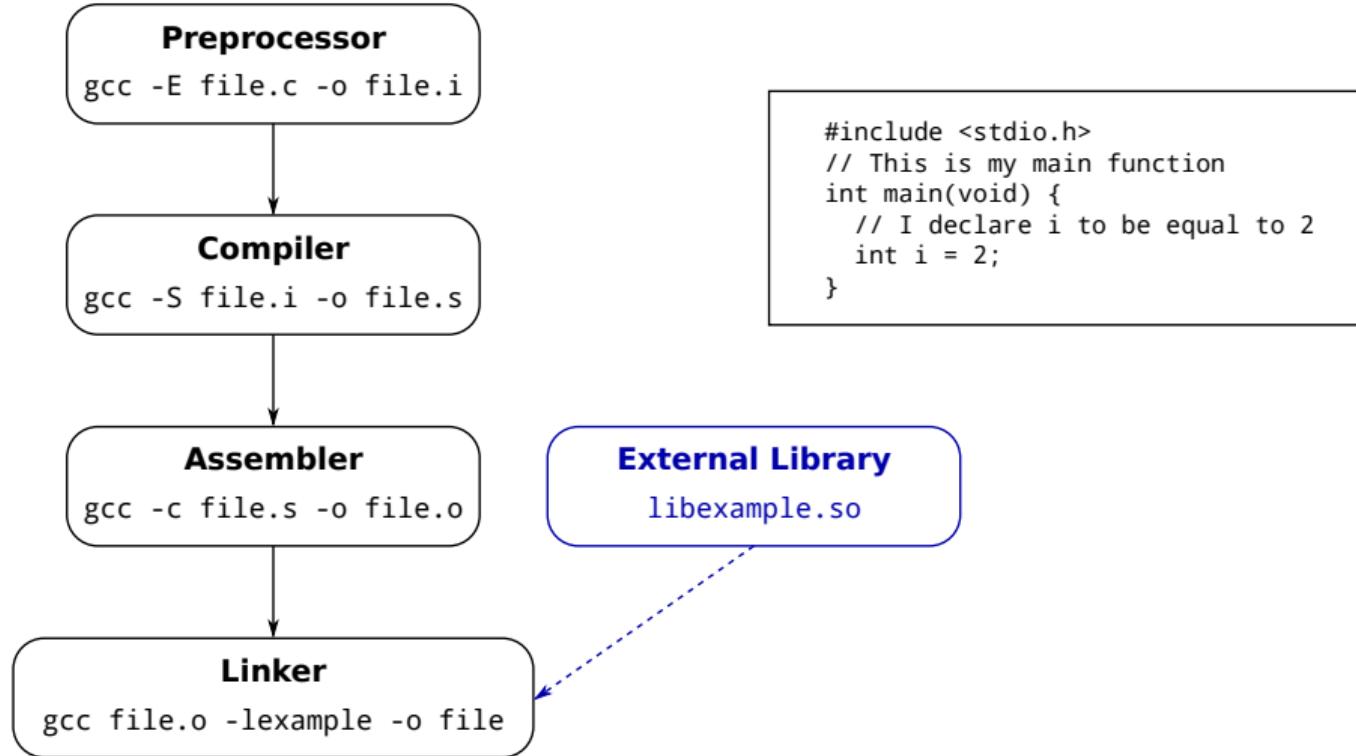
Assembly Translate the assembly into machine code and store it in object files

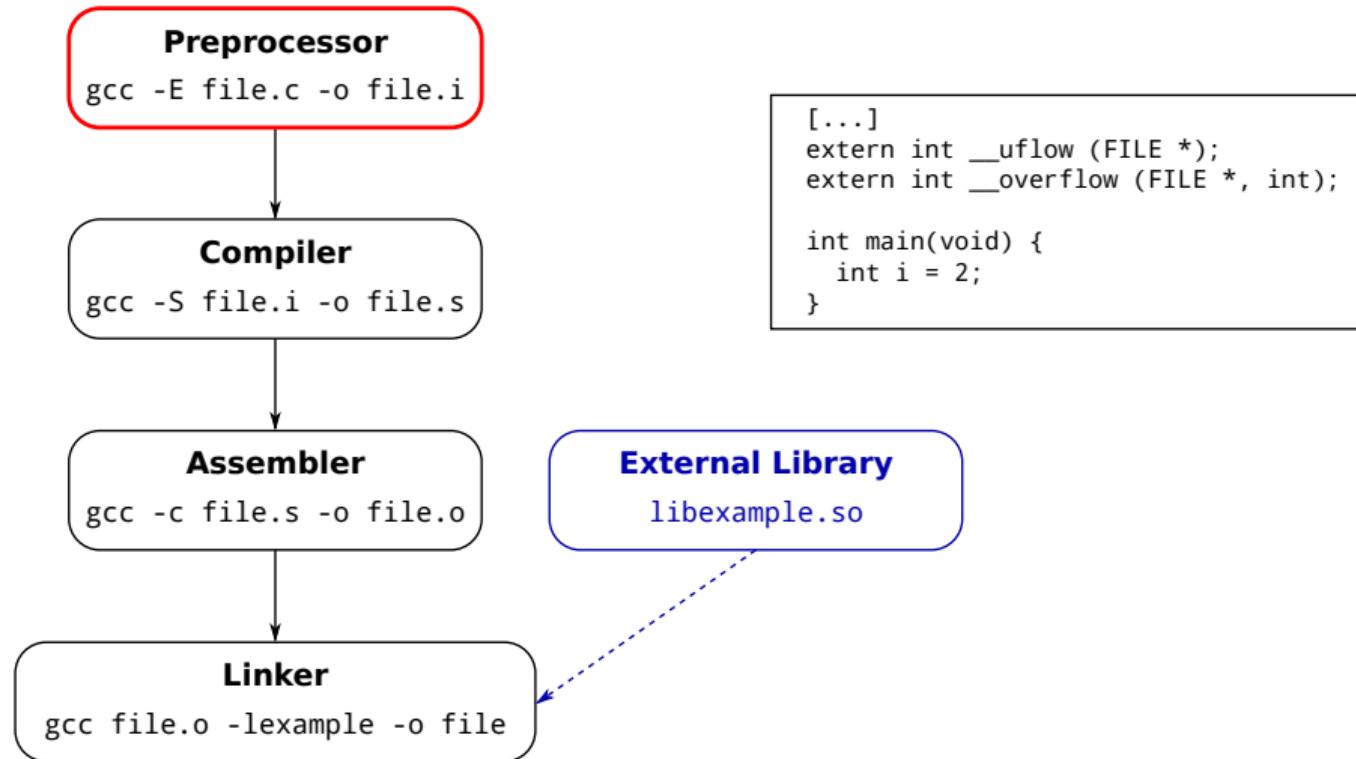
Linking Link all the object files into one executable

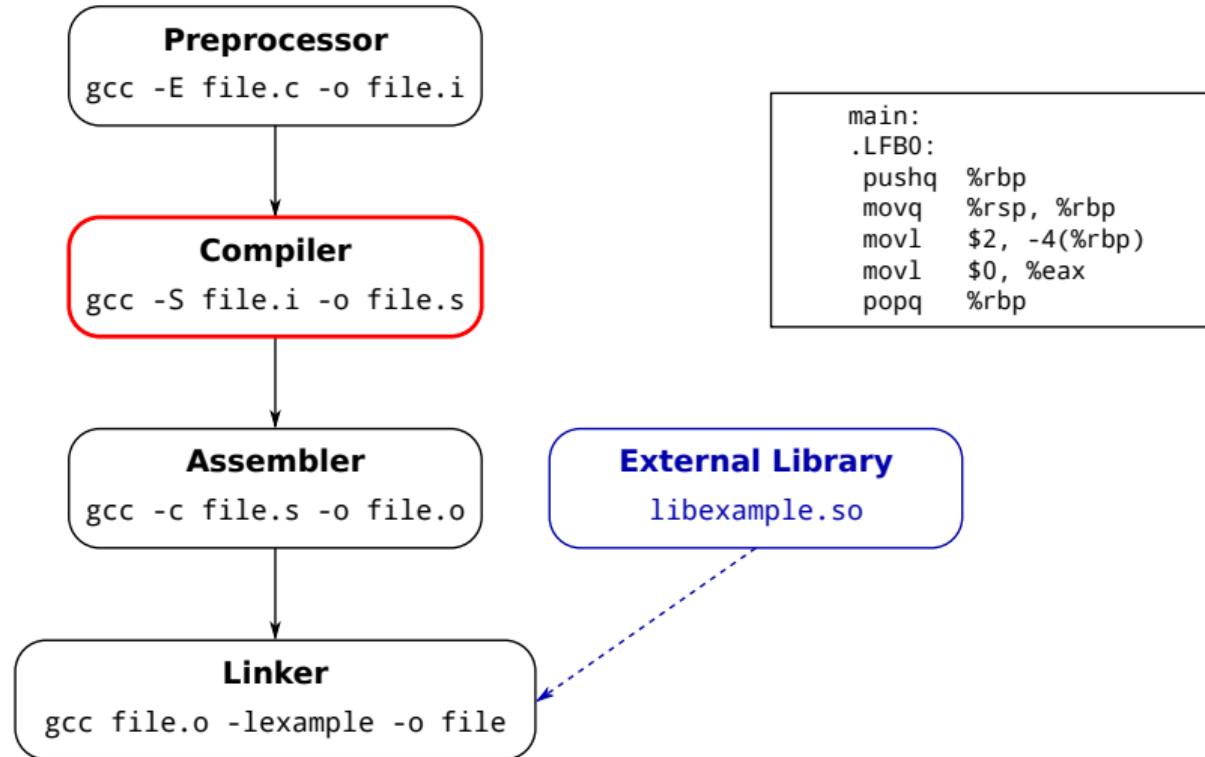
- In practice, the first three steps are combined together and simply called “compiling”

Compilation stages

The four compilation steps (visually)

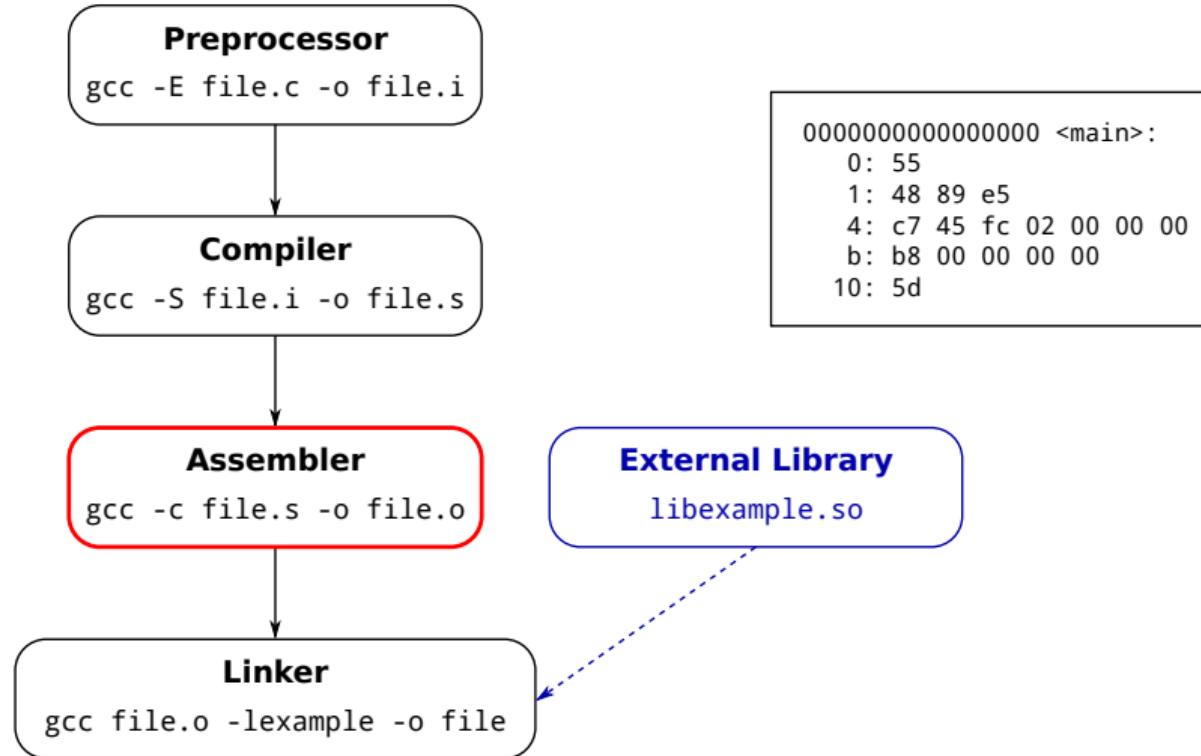






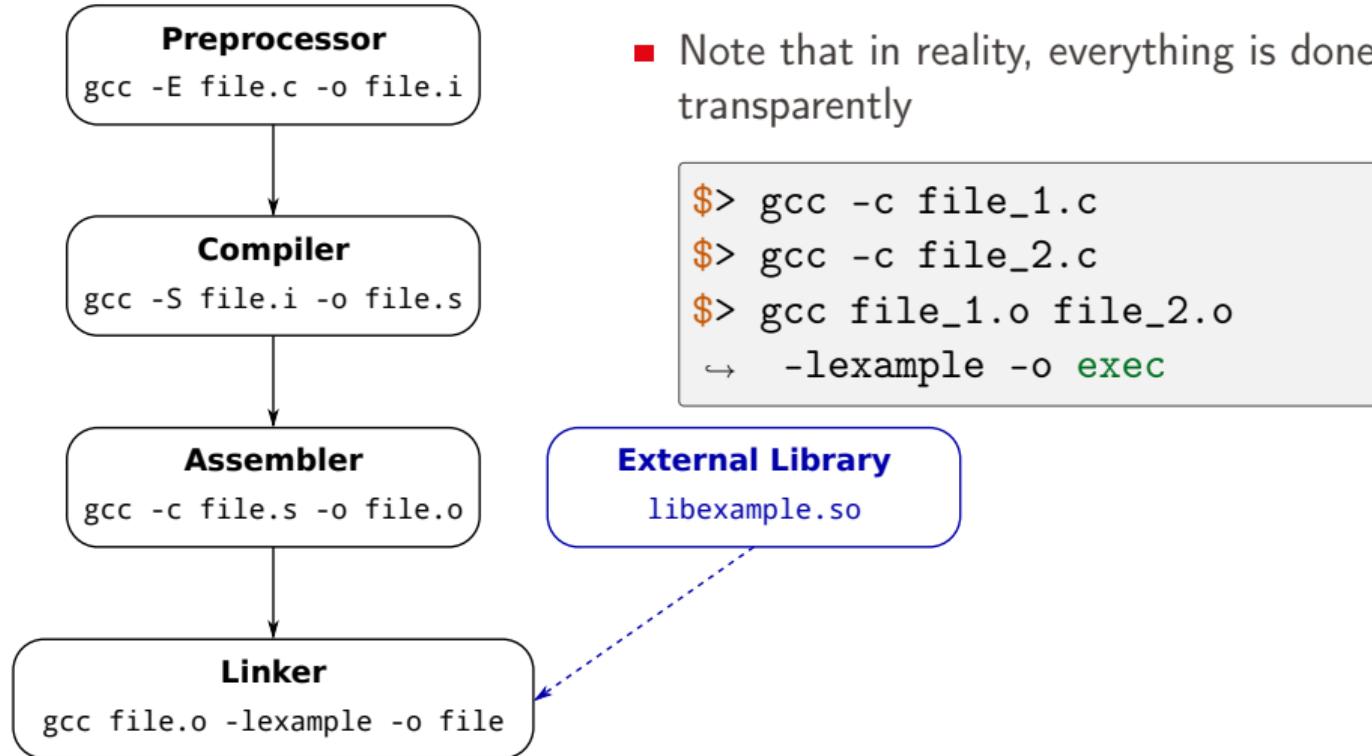
Compilation stages

The four compilation steps (visually)



Compilation stages

The four compilation steps (visually)



- Clone the repository

<https://gitlab.epfl.ch/math454-phpc/exercises-2025>

(you have to use your gaspar credentials)

```
$ git clone  
→ https://gitlab.epfl.ch/math454-phpc/exercises-2025.git
```

You can clone it either in your local machine or in helvetios.

- ▶ If you clone it in your local machine, you have to copy the folders to helvetios using scp.

- From helvetios, enter the folder **hello-separated-files**
- Generate `hello.o` and `greetings.o`
- Generate `hello` from `hello.o greetings.o`

What is makefile

- It is a build automation system
- It is a set of rules on how to produce an executable:
 - ▶ what and how to compile?
 - ▶ what and how to link?
- Rule for compiling

```
.cc.o:
    $(CXX) $(CPPFLAGS) $(CXXFLAGS) -c $<
```

- Rule for linking

```
$(EXE): $(OBJS)
    $(LD) -o $@ $(LDFLAGS) $(OBJS)
```

- Usage: `make`

- Enter the folder **hello-makefile**
- Read the **Makefile**
- Add the **-Wall -Werror** option to the compilation options
- Compile the code
- Modify the file **greetings.cc** as:

greetings.cc

```
1 #include <iostream>
2 #include "greetings.hh"
3
4 void greetings() {
5     int dummy;
6     std::cout << "Hello world !" << std::endl;
7 }
```

- Two types of libraries **static** or **shared**
 - ▶ static: archive of object file
 - ▶ shared: library loaded dynamically at execution

- Generate a shared library

Object files have to be compiled with **-fPIC** option

```
$ g++ -o libgreetings.so -shared -fPIC greetings.o
```

- Link

```
$ g++ -L<lib path> -lgreetings -o <exec> <obj_1> ... <obj_n>
```

- Execute

```
$ LD_LIBRARY_PATH=$LD_LIBRARY_PATH:<lib path> <exec>
```

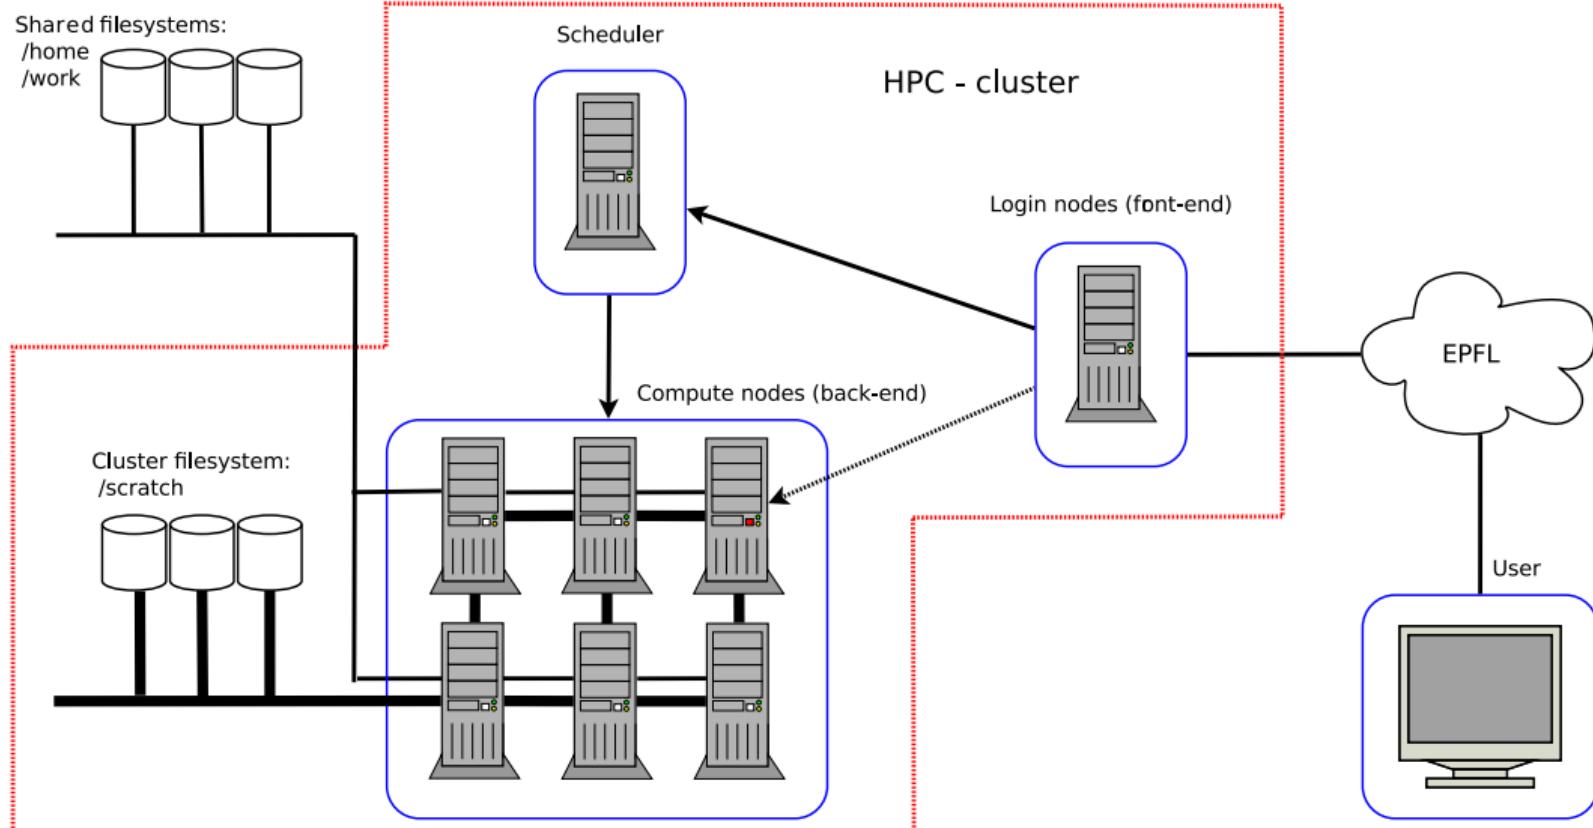
Real life project

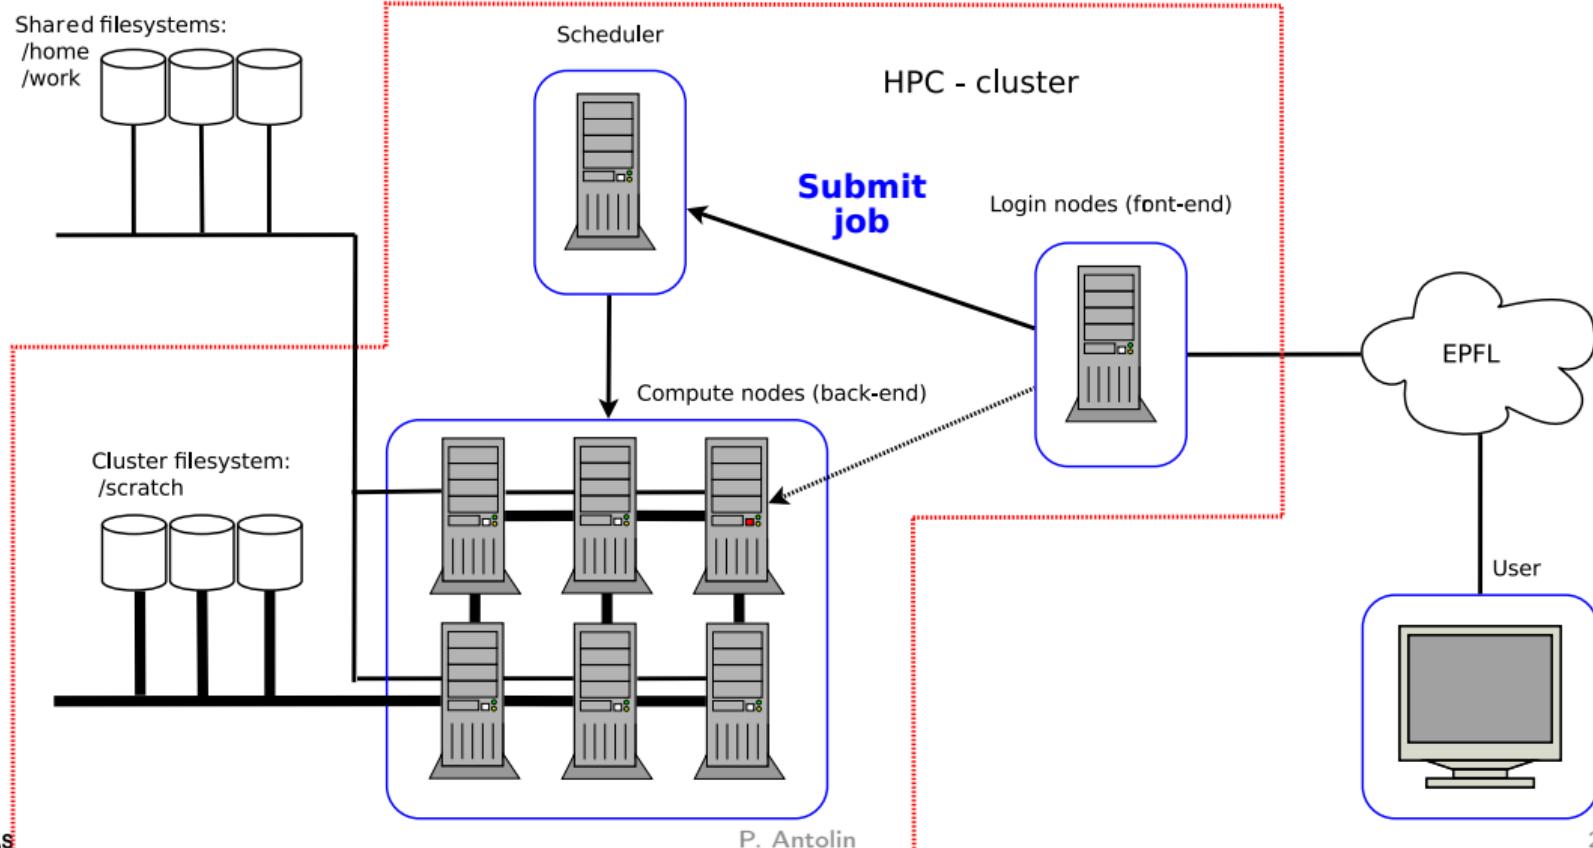
For project a bit more consequent with multiple dependencies to external libraries,
usage of a build automation system

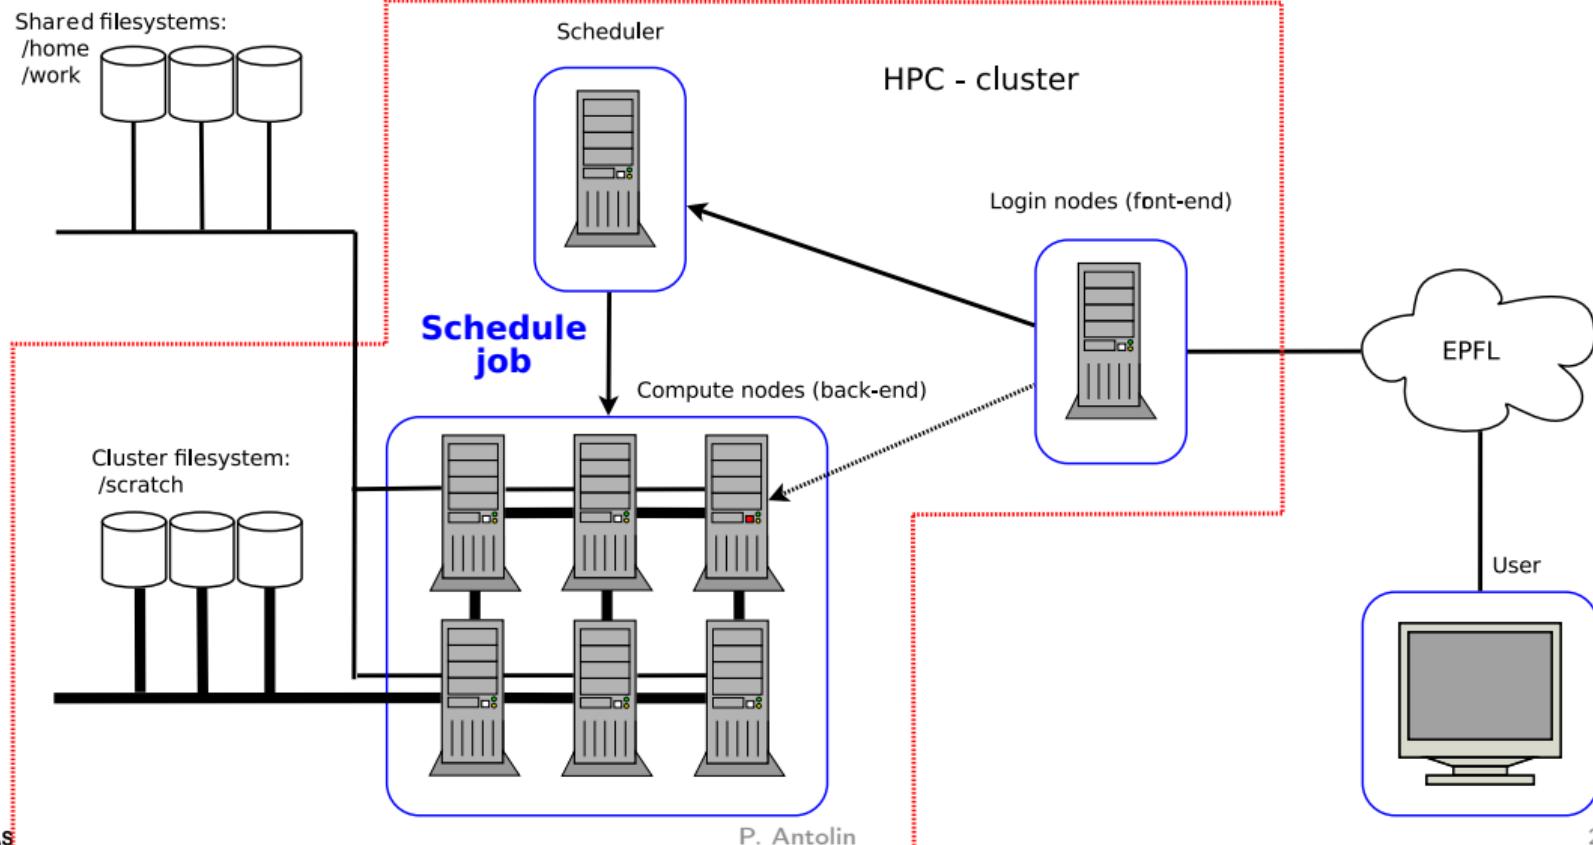
- CMake
- Autotools: configure
- SCons

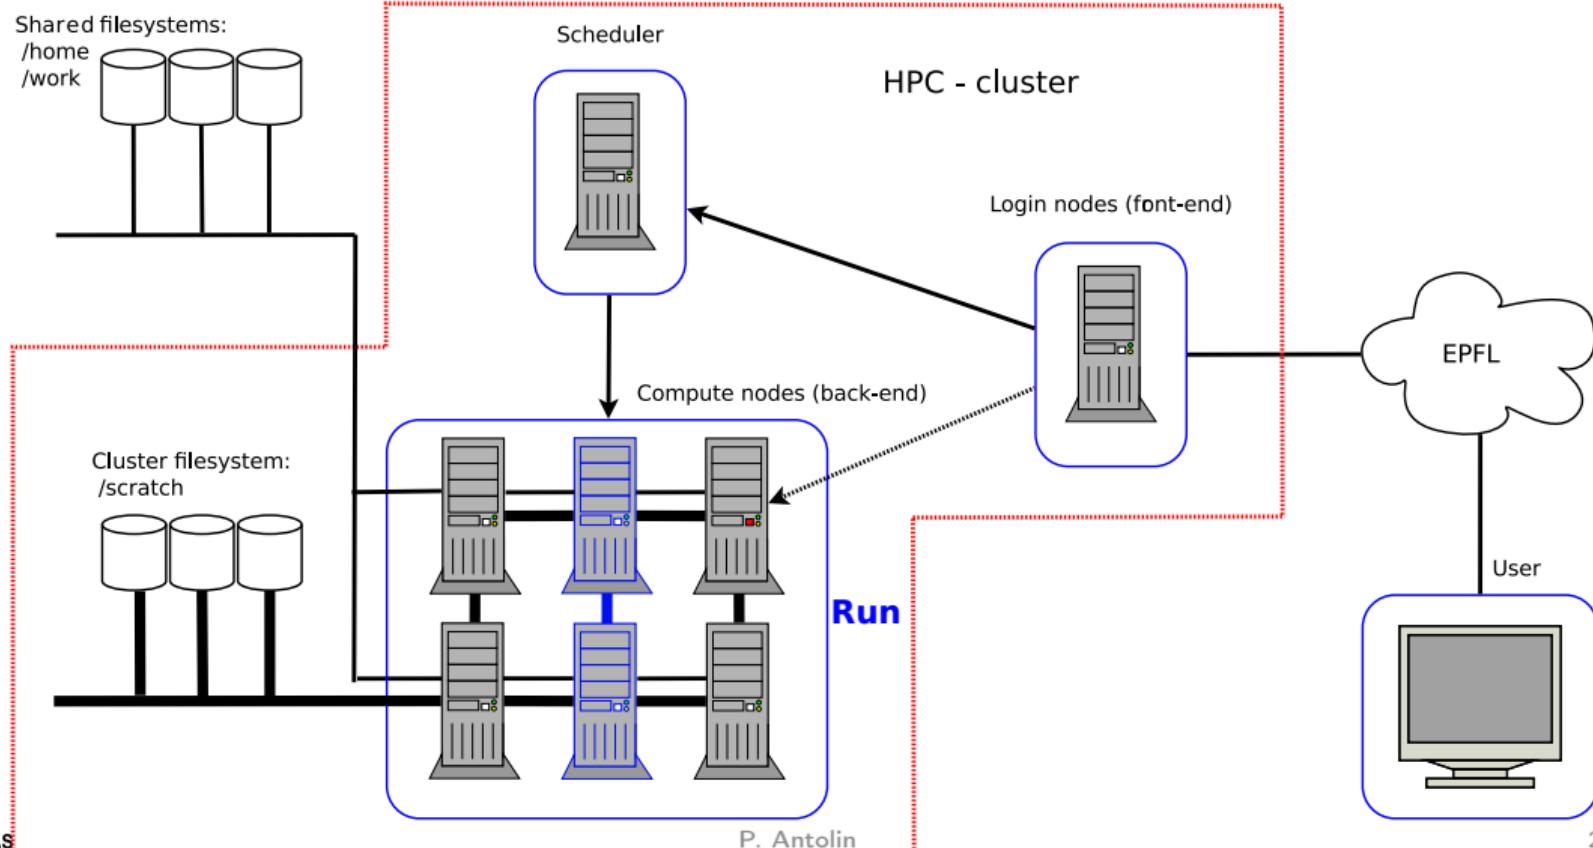












What is SLURM

- Simple Linux Utility for Resource Management
- Job scheduler
- Soda in Futurama

Basic commands

- **sbatch** submit a job to the queue
- **salloc** allocates resources
- **squeue** visualize the state of the queue

SLURM options

- **-q / --qos=<QOS>** sets the quality of service (affects the scheduling priority).
During the course set it to **--qos=math-454**
- **-t / --time=<HH:MM:SS>** set a limit on the total run time of the job
- **-N / --nodes=<N>** request that a minimum of **N** nodes be allocated to the job
- **-n / --tasks=<n>** advises SLURM that this job will launch a maximum of **n**, in the MPI sense
- **-c / --cpus-per-task=<ncpus>** advises SLURM that job will require **ncpus** per task
- **--ntasks-per-node=<ntasks>** number of tasks per node
- **--mem=<size[units]>** defines the quantity of memory per node requested

Need more help? Have a look at <https://scitas-doc.epfl.ch> and

<https://slurm.schedmd.com/sbatch.html>

Or you can put everything in a file, e.g.

mysimulation.job

```
#!/bin/bash -l
#SBATCH --qos=math-454
#SBATCH --time=01:10:00
#SBATCH --nodes=2
#SBATCH --ntasks=72

./myprogram
```

and submit the job

```
$ sbatch mysimulation.job
```

To list all your running jobs

```
$ squeue -u <username>
```

To cancel a simulation

```
$ scancel <jobid>
```

The <jobid> can be found using squeue

To allocate a node

```
$ salloc --qos=math-454 --account=math-454
```

- Check the queue state
- Allocate one node with QOS math-454
- Try **srun hostname**, we will see this command more in detail in later exercise sessions
- Exit the allocation to not block resources: **exit** or **Ctrl-d** (**IMPORTANT !!**)

- Write a script that runs the hello world code

```
myscript.job
```

```
#!/bin/bash -l  
./hello
```

- Try your script.

note: in general you should not try your codes on the front node

```
$ sh myscript.job
```

- Submit your script with **sbatch** and adding the QOS

```
$ sbatch --qos=math-454 myscript.job
```

- Try

```
$ squeue -u <username>
```

- A file named slurm-<jobid>.out should have been created, check its content

- Add the QOS as an option directly in the script

```
myscript.job
```

```
#!/bin/bash -l
#SBATCH --qos=math-454
./hello
```

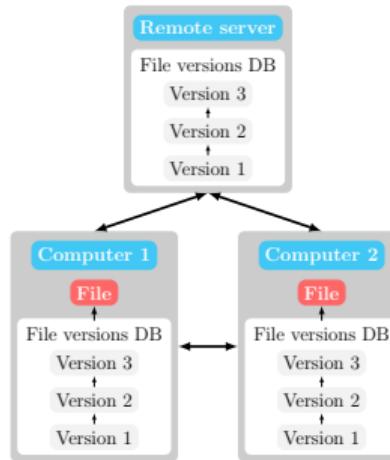
- Submit it again

```
$ sbatch myscript.job
```



Git: *the stupid content tracker*

- Distributed revision control
- Originally developed by Linus Torvald
- Named after the *egotistical bastard* Linus



MATH-454 Parallel and High Performance Computing Lecture

1: Execution in an HPC environment

- Basics on GIT

- “Versioning”: with Git

Git means “unpleasant person”, “I’m an egotistical bastard, and I name all my projects after myself. First ‘Linux’, now ‘git’.” – Linus



“Versioning”: with Git

Git: *the stupid content tracker*

- Distributed revision control
- Originally developed by Linus Torvald
- Named after the egotistical bastard Linus



REMOTE SERVER



```
$ git clone <uri repo.git>
```



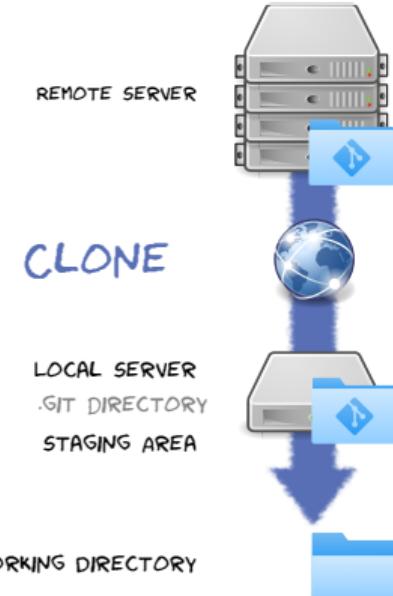
LOCAL SERVER
.GIT DIRECTORY
STAGING AREA



WORKING DIRECTORY



```
$ git clone <uri repo.git>
Cloning into '<repo>'...
remote: Counting objects: 6940, done.
remote: Total 6940 (delta 0), reused ...
Receiving objects: 100% (6940/6940), ...
Resolving deltas: 100% (3286/3286), done.
```



```
$ git clone <uri repo.git>
Cloning into '<repo>'...
remote: Counting objects: 6940, done.
remote: Total 6940 (delta 0), reused ...
Receiving objects: 100% (6940/6940), ...
Resolving deltas: 100% (3286/3286), done.
```

```
$ git status
```

```
$ git status
```

```
On branch master
```

```
Your branch is up-to-date with 'origin/master'.
```

```
nothing to commit, working tree clean
```



STAGING AREA



WORKING DIRECTORY

```
$ git status
On branch master
Your branch is up-to-date with 'origin/master'.

Untracked files:
  (use "git add <file>..." to include in what will be committed)

    my_code.py

nothing added to commit but untracked files present
```



```
$ git status
On branch master
Your branch is up-to-date with 'origin/master'.

Changes to be committed:
(use "git reset HEAD <file>..." to unstage)

    new file:   my_code.py
```

REMOTE SERVER



COMMIT



LOCAL SERVER

.GIT DIRECTORY
STAGING AREA

WORKING DIRECTORY



```
$ git commit -m <message>
```

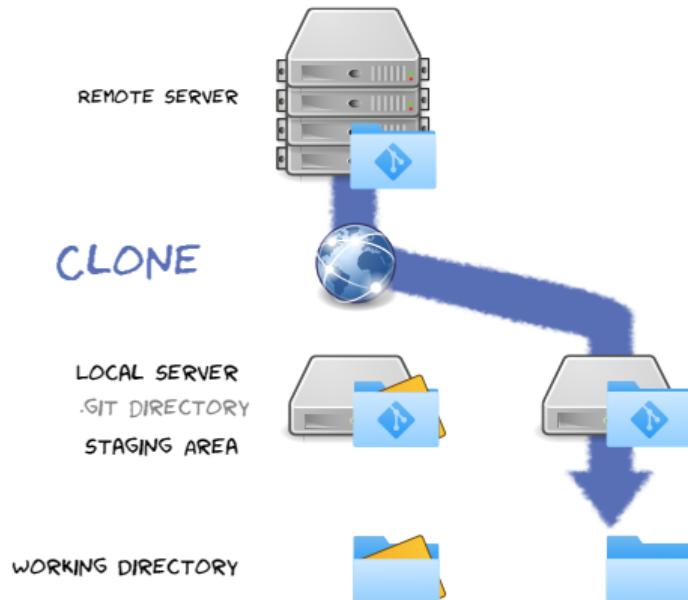
```
$ git status
```

```
On branch master
```

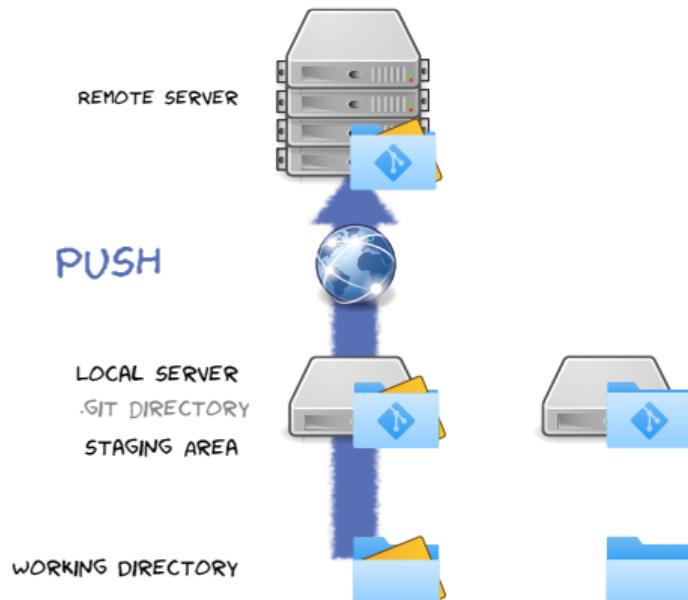
```
Your branch is ahead of 'origin/master' by 1 commit.
```

```
(use "git push" to publish your local commits)
```

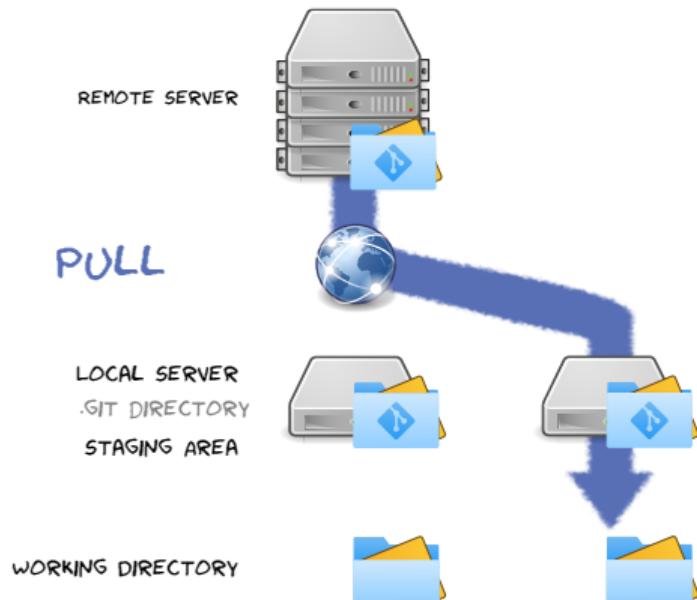
```
nothing to commit, working tree clean
```



```
$ git clone <uri>
```



```
$ git push
```

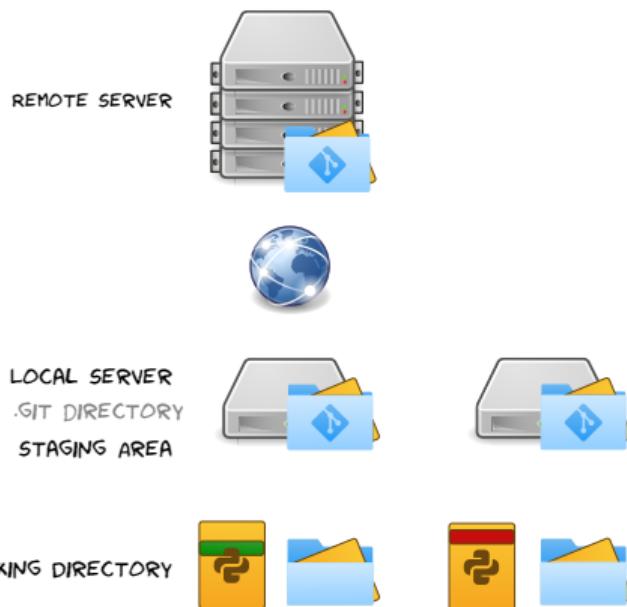


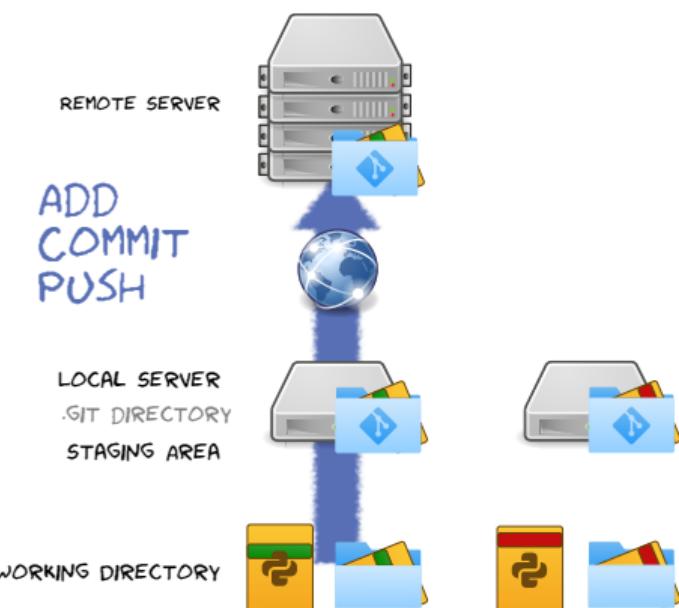
```
$ git pull
```

- If you do not have git installed, get it from <https://git-scm.com/downloads> or from your package manager
- Go on <https://gitlab.epfl.ch/> and login with your EPFL account.
- Once connected go on your user settings page (circle icon top left corner)
- In the **Preferences > SSH Keys > Add new key** menu set your public ssh key. This key will be used to connect to the git server through ssh.

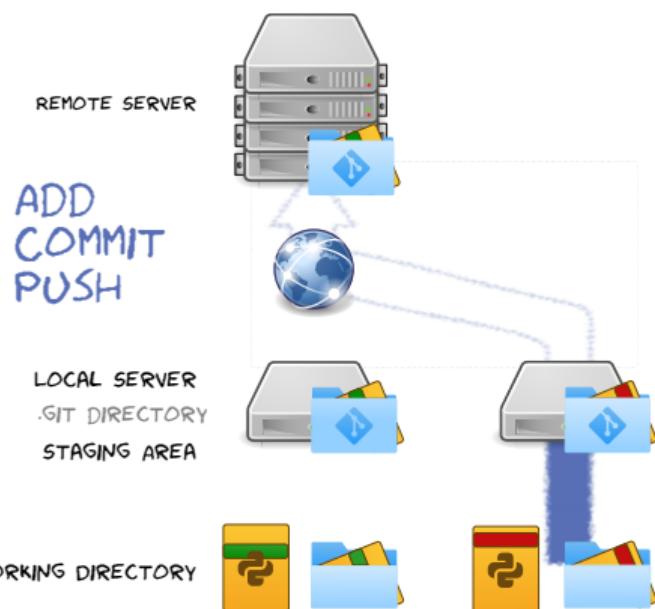
git clone <repo url> [local name]	Clone a remote repository
git add <files...>	Stage modified/new files
git commit -m "comment"	Commit staged files
git pull	Pull and merge remote modifications
git push	Push the local modifications to the remote server
git status	Check the local state

- Now you should be able to clone a repository
Either create a repository or clone <git@gitlab.epfl.ch:math454-phpc/test-repo.git>
- Create a file, use a filename that will not clash with the others
- Check the state of your working copy



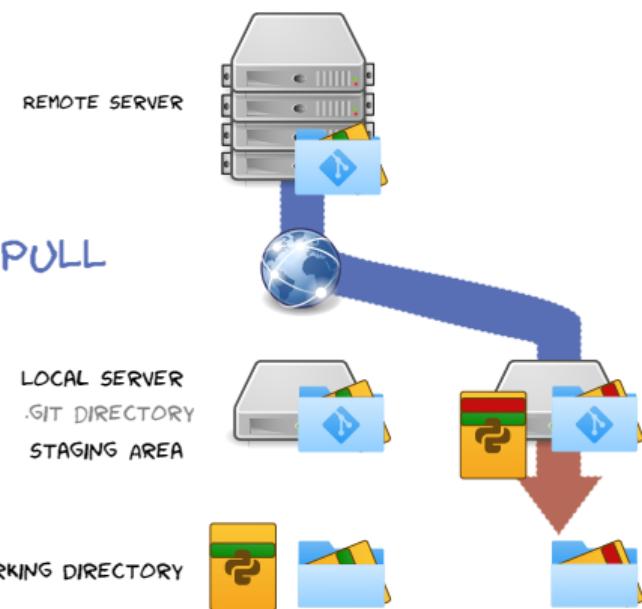


```
$ git add <filename>
$ git commit -m <message>
$ git push
```



```
$ git add <filename>
$ git commit -m <message>
$ git push
```

```
$ git push
To <repo>
! [rejected]           master -> master (fetch first)
error: failed to push some refs to '<repo>'
hint: ...
```



```
$ git pull
```

```
$ git pull
remote: Counting objects: 3, done.
remote: Total 3 (delta 0), reused 0 (delta 0)
Unpacking objects: 100% (3/3), done.
From <repo>
fe22d81..0bcfb99  master      -> origin/master
Auto-merging my_code.py
CONFLICT (content): Merge conflict in my_code.py
Automatic merge failed; fix conflicts and then commit the result.
```

```
$ git status
```

On branch master

Your branch and 'origin/master' have diverged,
and have 1 and 1 different commits each, respectively.
(use "git pull" to merge the remote branch into yours)

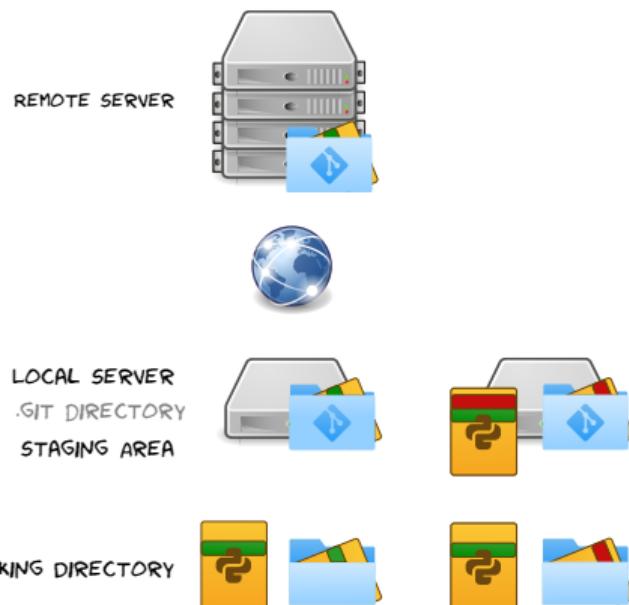
You have unmerged paths.

(fix conflicts and run "git commit")

(use "git merge --abort" to abort the merge)

Unmerged paths:

(use "git add <file>..." to mark resolution)



Correct the conflict:

my_file.py

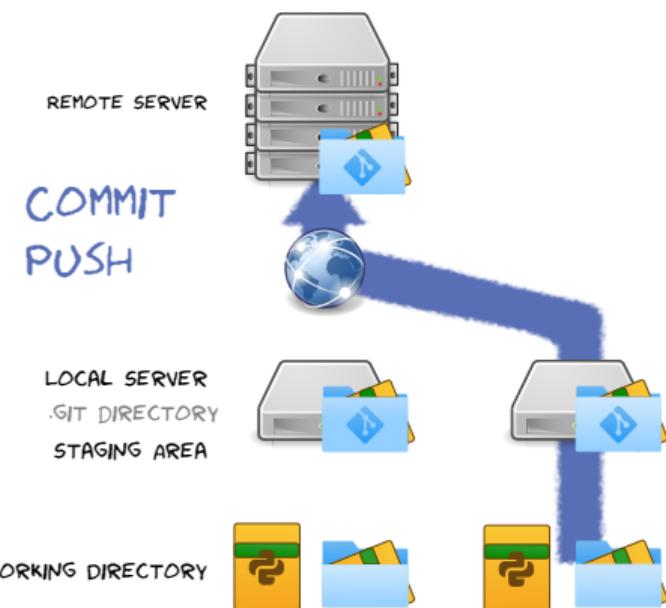
<<<<<<

One version

=====

Other version

>>>>>>



my_file.py

One version

```
$ git commit -a  
$ git push
```

- Modify the file created in the previous exercise in both clones
- Commit this both modifications
- Pull and push in one of the clone
- Pull in the second clone, You should get a conflict

```
<<<<<<<
```

```
One version
```

```
=====
```

```
Other version
```

```
>>>>>>>
```

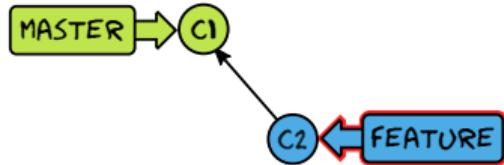
- Check the local status
- Correct the conflict and commit using **git commit -a**
- Push the modifications



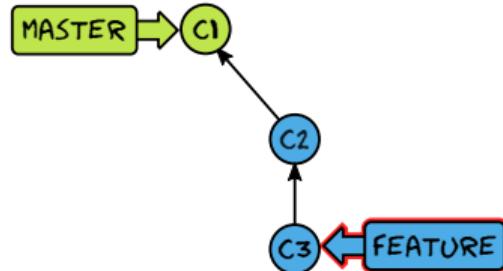
```
$ git clone <uri repo.git>
```



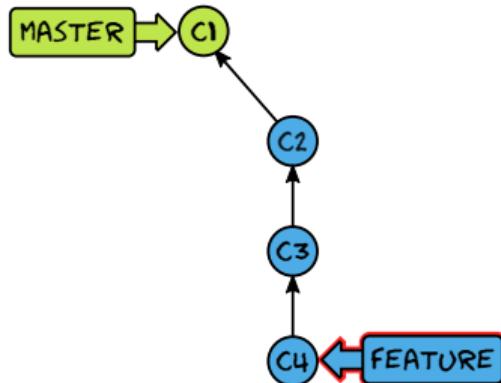
```
$ git checkout -b feature
```



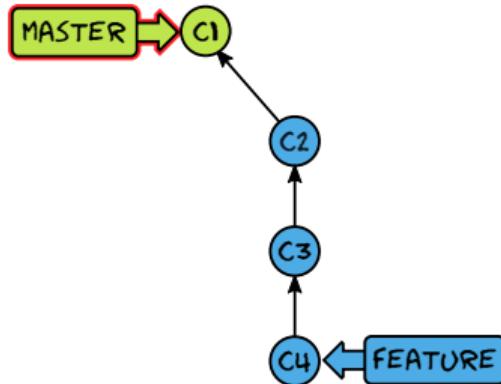
```
$ git commit -m <message>
```



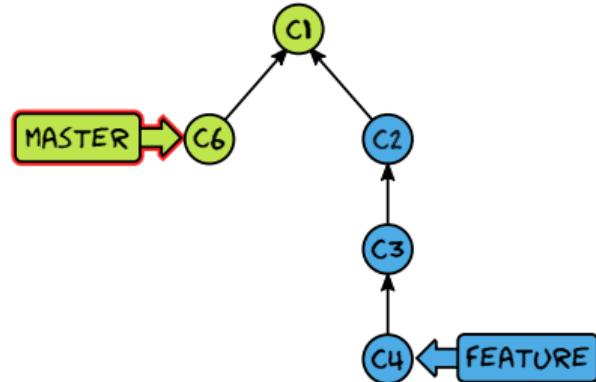
```
$ git commit -m <message>
```



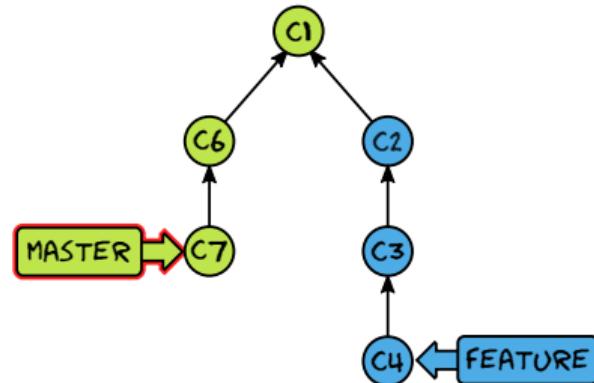
```
$ git commit -m <message>
```



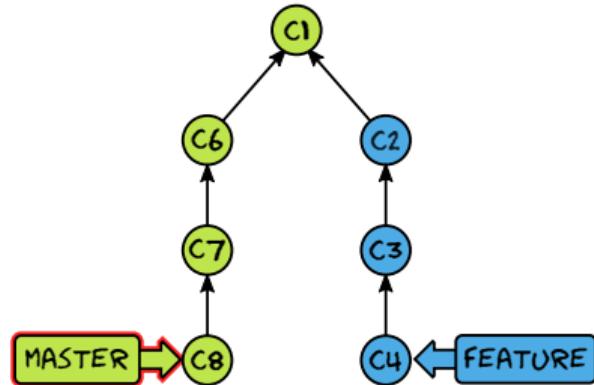
```
$ git checkout master
```



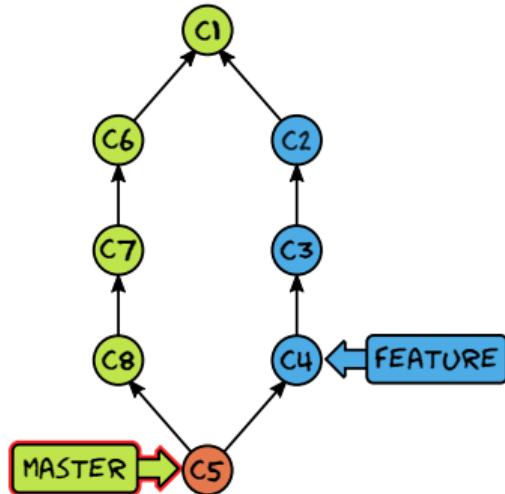
```
$ git commit -m <message>
```



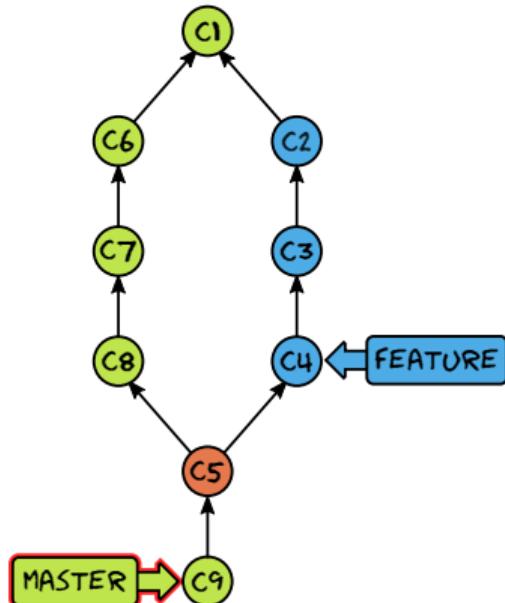
```
$ git commit -m <message>
```



```
$ git commit -m <message>
```

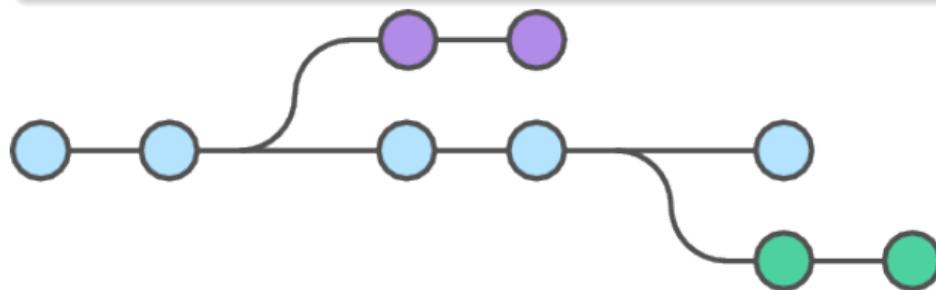


```
$ git merge feature
```

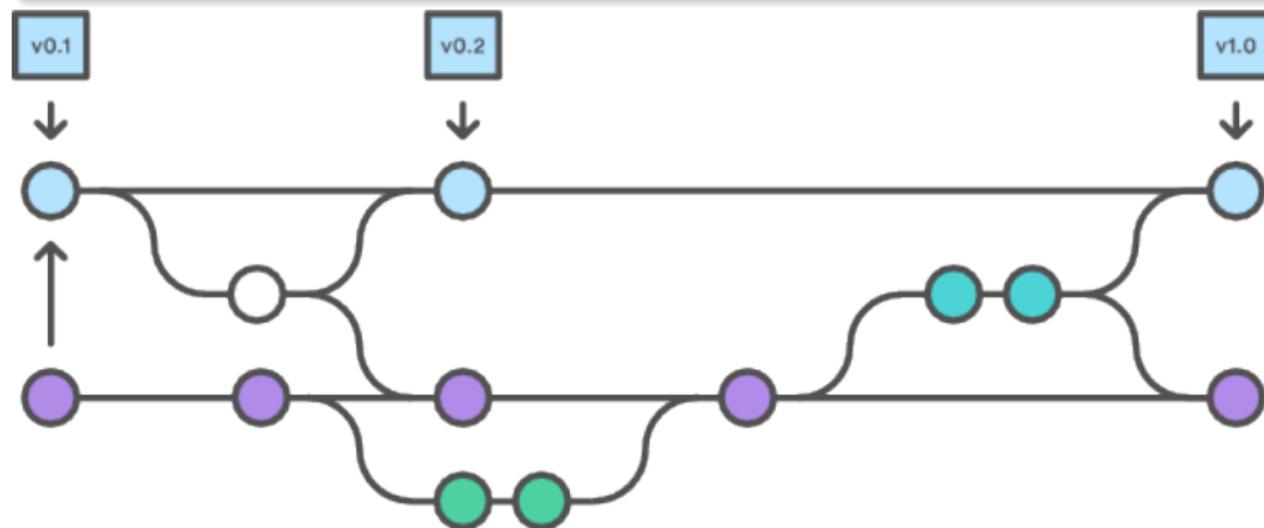


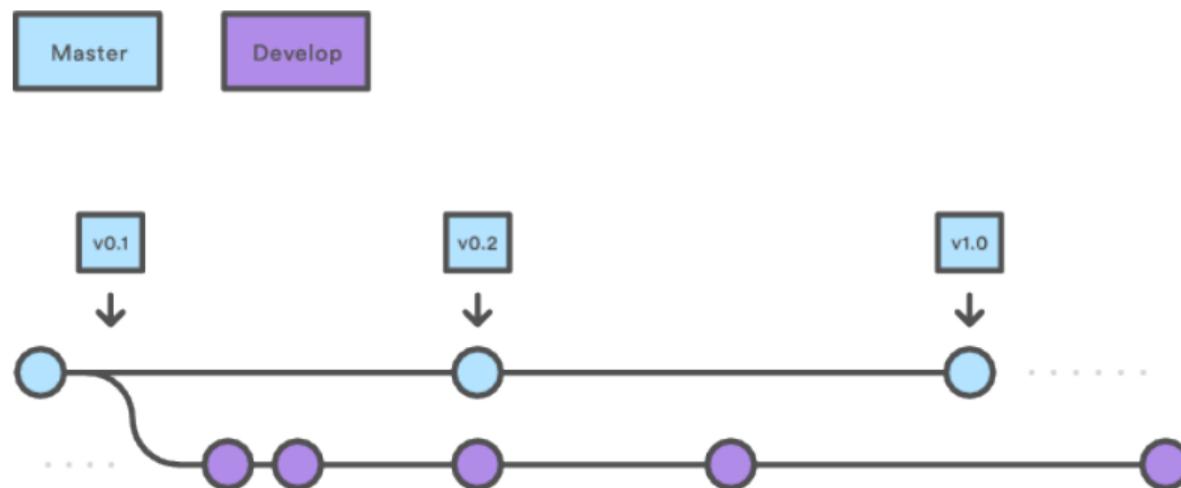
```
$ git commit -m <message>
```

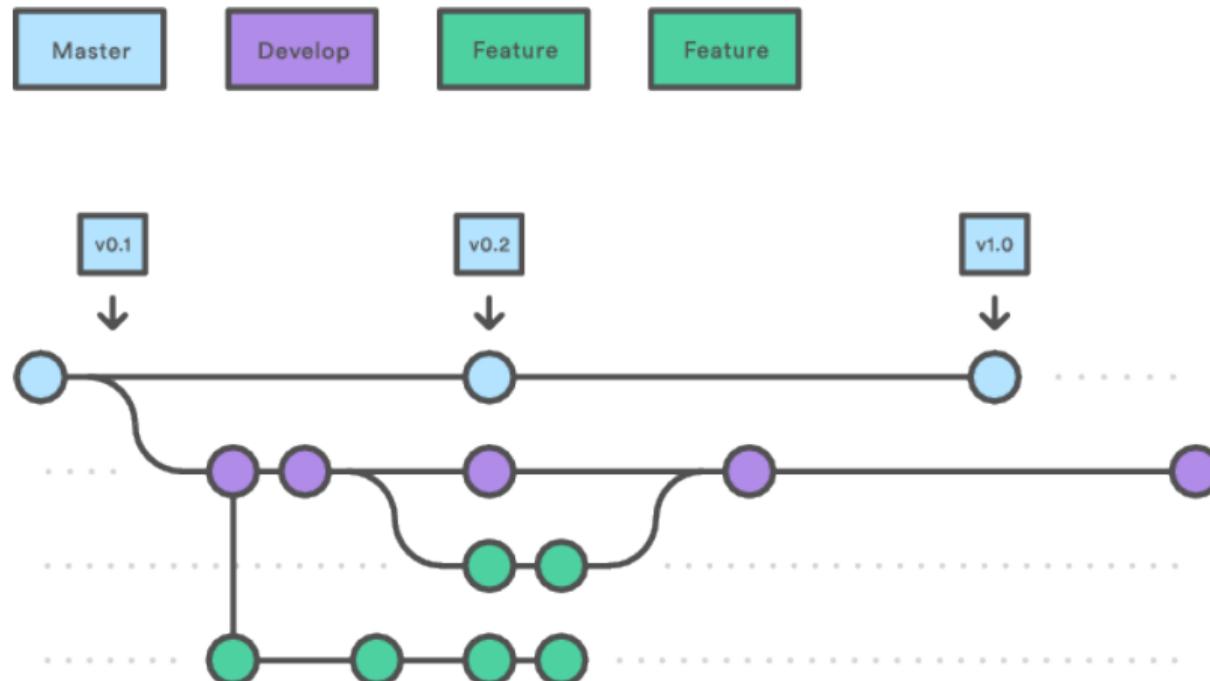
Feature branch

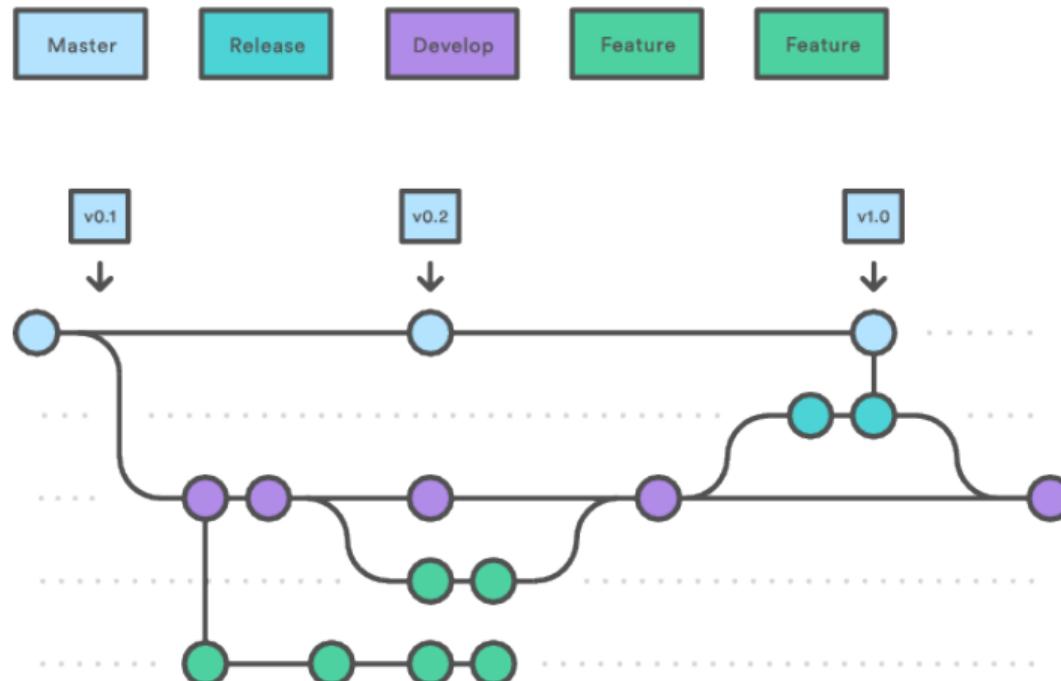


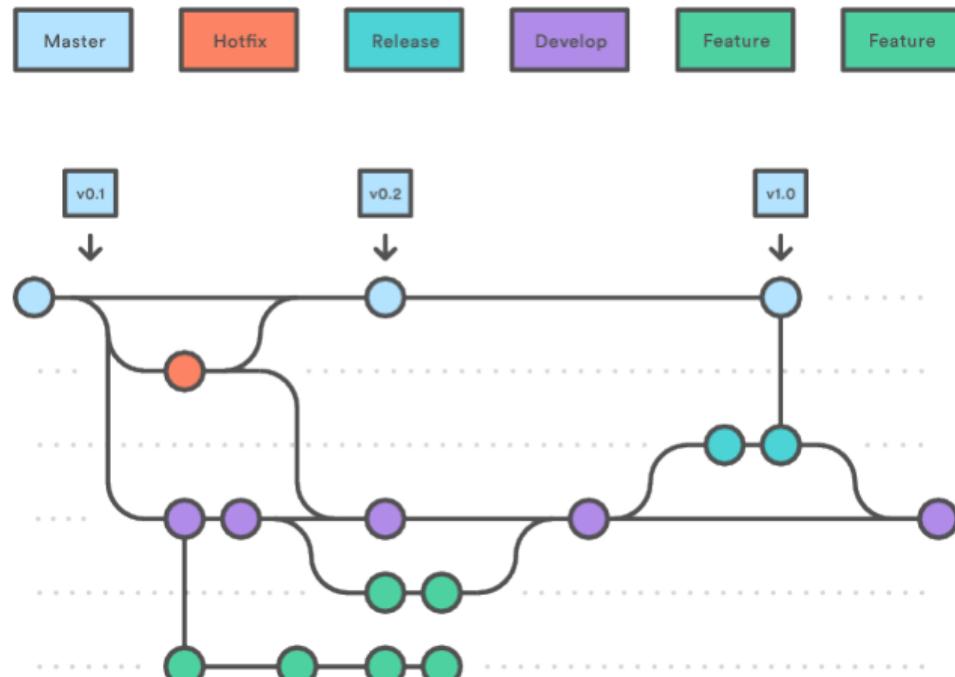
Gitflow











git branch <name>	Create a new branch from the current HEAD
git checkout <name>	Switch to the specified branch
git merge <name>	Merge the branch specified in the current one
git branch -d	Delete a branch
git branch -a	List all branches
git log	List the different commits of the current branch
git log --graph --all	Show also the branches

- Create a branch with the name of your choice
- Modify a file and commit the changes
- Checkout the **master** branch
- Modify a file and commit the changes
- Merge the branch previously created in the **master** branch
- List all branches
- Print the logs of the different modifications

Sources

- Wikipedia
- <http://git-scm.com>
- Manpages: rsync, git
- <https://www.atlassian.com/git/>
- <http://nvie.com/posts/a-successful-git-branching-model/>

Learn more

- Git with a game: <http://learngitbranching.js.org/>