


# MATH-454 Parallel and High Performance Computing

## Lecture 4: MPI basics

Pablo Antolin

Slides of N. Richart, E. Lanti, V. Keller's lecture notes


March 20 2025




# Message Passing Interface (MPI)



- Introduce distributed memory programming paradigm.
- Point-to-point communications.
- Collective communications.

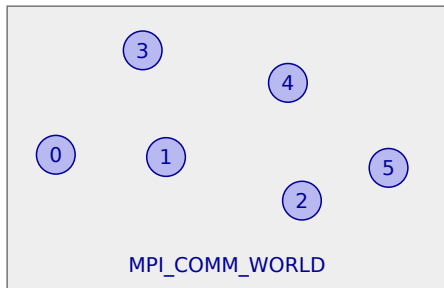




- MPI is a *Message-Passing Interface* specification.
- There are many implementations (Intel MPI, OpenMPI, MPICH, MVAPICH, etc.)
- Library interface, not a programming language.
- It is standardized:
  - ▶ Defined by the [MPI forum](#)
  - ▶ Current version is MPI 4.1
- As such, it is portable, flexible, and efficient.
- Interface to C and Fortran in standard. C interface usable with C++. Also MPI bindings for other languages (Python, Julia, MATLAB, etc.).

## mpi/hello\_mpi.cc

```
1  #include <iostream>
2  #include <mpi.h>
3
4  int main(int argc, char *argv[]) {
5      MPI_Init(&argc, &argv);
6
7      int size, rank;
8      MPI_Comm_size(MPI_COMM_WORLD, &size);
9      MPI_Comm_rank(MPI_COMM_WORLD, &rank);
10
11     std::cout << "I am process " << rank << " out of " << size <<
        ↪     std::endl;
12
13     MPI_Finalize();
14     return 0;
15 }
```



- MPI code is bounded by a `MPI_Init` and a `MPI_Finalize`.
- MPI starts  $N$  processes numbered  $0, 1, \dots, N - 1$ .  
The number of every process is usually denoted as *rank*.
- They are grouped in a *communicator* of size  $N$ .
- After init, MPI provides a default communicator called `MPI_COMM_WORLD`.




- Point-to-Point (One-to-One).
- Collectives (One-to-All, All-to-One, All-to-All).
- One-sided/Shared memory (One-to ...).
- Blocking and Non-Blocking of all types.

## Syntax

```
1 int MPI_Ssend(const void *buf, int count, MPI_Datatype datatype,  
  ↪ int dest, int tag, MPI_Comm comm);  
2  
3 int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int  
  ↪ source, int tag, MPI_Comm comm, MPI_Status *status);
```

- **buf**: pointer to the data to send/receive.
- **count**: number of elements to send/receive.
- **datatype**: datatype of the data to send/receive.
- **dest, source**: the rank of the destination/source of the communication.
- **tag**: a message tag to differentiate the communications.
- **comm**: communicator in which the communication happens.
- **status**: object containing information on the communication.



- Buffer is a pointer to the first data (buf), a size (count), and a datatype.
- Datatypes (extract):
  - ▶ MPI\_INT
  - ▶ MPI\_UNSIGNED
  - ▶ MPI\_FLOAT
  - ▶ MPI\_DOUBLE
- For `std::vector<double> vect`:
  - ▶ `buf = vect.data()`
  - ▶ `count = vect.size()`
  - ▶ `datatype = MPI_DOUBLE`

## Constants:

- `MPI_STATUS_IGNORE`: to state that the status is ignored.
- `MPI_PROC_NULL`: placeholder for the source or destination.
- `MPI_ANY_SOURCE`: is a wildcard for the source of a receive.
- `MPI_ANY_TAG`: is a wildcard for the tag of a receive.

`MPI_Status`:

- Structure containing tag and source

```
1 MPI_Status status;  
2 std::cout << "Tag: " << status.tag << " - Source: " <<  
  ↪ status.source << std::endl;
```

- Size of the received buffer can be asked using the status:

```
1 int MPI_Get_count(const MPI_Status *status, MPI_Datatype  
  ↪ datatype, int *count);
```

## mpi/send\_recv.cc

```
16 MPI_Init(NULL, NULL);
17 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18 MPI_Comm_size(MPI_COMM_WORLD, &size);
19
20 assert(size == 2 && "Works only with 2 procs");
21
22 int tag = 0;
23
24 if (rank == 0) {
25     fill_buffer(buf);
26     MPI_Ssend(buf.data(), buf.size(), MPI_INT, 1, tag, MPI_COMM_WORLD);
27 } else if (rank == 1) {
28     MPI_Recv(buf.data(), buf.size(), MPI_INT, 0, tag, MPI_COMM_WORLD,
29             ↪ MPI_STATUS_IGNORE);
29 }
```

<u>Rank 0</u>		<u>Rank 1</u>
send(1,0)	→	recv(0,0)
finalize()		finalize()

- `MPI_Ssend`: (S for Synchronous) function returns when other end posted matching `recv` and the buffer can be safely reused.
- `MPI_Bsend`: (B for Buffer) function returns immediately, send buffer can be reused immediately.
- `MPI_Rsend`: (R for Ready) can be used only when a receive is already listening (waiting).
- `MPI_Send`: acts like `MPI_Bsend` on small arrays, and like `MPI_Ssend` on bigger ones.

- `MPI_Ssend`: (S for Synchronous) function returns when other end posted matching `recv` and the buffer can be safely reused.
- `MPI_Bsend`: (B for Buffer) function returns immediately, send buffer can be reused immediately.
- `MPI_Rsend`: (R for Ready) can be used only when a receive is already listening (waiting).
- `MPI_Send`: acts like `MPI_Bsend` on small arrays, and like `MPI_Ssend` on bigger ones.

Be careful with deadlocks and race conditions!!



## mpi/send\_recv\_deadlock.cc

```
25 assert(size == 2 && "Works only with 2 procs");
26 int tag = 0;
27
28 if (rank == 0)
29 {
30     fill_buffer(buf);
31     MPI_Ssend(buf.data(), buf.size(), MPI_INT, 1, tag, MPI_COMM_WORLD);
32     MPI_Recv(buf.data(), buf.size(), MPI_INT, 1, tag, MPI_COMM_WORLD,
33             ↪ MPI_STATUS_IGNORE);
34 }
35 else if (rank == 1)
36 {
37     MPI_Ssend(buf.data(), buf.size(), MPI_INT, 0, tag, MPI_COMM_WORLD);
38     MPI_Recv(buf.data(), buf.size(), MPI_INT, 0, tag, MPI_COMM_WORLD,
39             ↪ MPI_STATUS_IGNORE);
40 }
```

<u>Rank 0</u>	<u>Rank 1</u>
send(1,0)	send(0,0)
barrier!	barrier!

## mpi/send\_recv\_deadlock\_2.cc

```
16 MPI_Init(NULL, NULL);
17 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18 MPI_Comm_size(MPI_COMM_WORLD, &size);
19
20 // Let's now assume that we may have 2 or more processes.
21 int tag = 0;
22
23 if (rank == 0) {
24     fill_buffer(buf);
25     MPI_Ssend(buf.data(), buf.size(), MPI_INT, 1, tag, MPI_COMM_WORLD);
26 } else {
27     MPI_Recv(buf.data(), buf.size(), MPI_INT, 0, tag, MPI_COMM_WORLD,
28             ↪ MPI_STATUS_IGNORE);
29 }
29 MPI_Finalize();
```

With 2 processes:

<u>Rank 0</u>		<u>Rank 1</u>
send(1,0)	→	recv(0,0)
finalize()		finalize()

With 3 processes:

<u>Rank 0</u>		<u>Rank 1</u>	<u>Rank 2</u>
send(1,0)	→	recv(0,0)	recv(0,0)
finalize()		finalize()	<b>barrier!</b>

## Syntax

```
1 int MPI_Sendrecv(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, int dest, int sendtag, void *recvbuf, int recvcount,  
  ↪ MPI_Datatype recvtype, int source, int recvtag, MPI_Comm comm,  
  ↪ MPI_Status *status);
```

- Combines a send and a receive, to help mitigate deadlocks.
- Has a in-place variant MPI\_Sendrecv\_replace.

## Syntax

```
1 int MPI_Isend(const void *buf, int count, MPI_Datatype datatype,  
  ↪ int dest, int tag, MPI_Comm comm, MPI_Request *request);  
2  
3 int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int  
  ↪ source, int tag, MPI_Comm comm, MPI_Request *request);
```

- I for *immediate*.
- request in addition to parameters from blocking version.  
It is an object attached to the communication.
- Receive does not have a status.
- The communication starts but is not completed.
- S, B, and R variants are also defined.

## Syntax

```
1 int MPI_Wait(MPI_Request *request, MPI_Status *status);  
2  
3 int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
```

- For non-blocking communications, completion should be “manually” checked.
- MPI\_Test or MPI\_Wait.
- Send completed means the buffer can be reused.
- Receive completed means the buffer can be read.
- status is set at completion.
- flag is **true** if completed, **false** otherwise.



## mpi/isend\_\_recv.cc

```
22 MPI_Request request;
23 if (rank == 0) {
24     fill_buffer(buf);
25     MPI_Isend(buf.data(), buf.size(), MPI_INT, 1, 0, MPI_COMM_WORLD,
26             ↪ &request);
27 } else if (rank == 1) {
28     MPI_Recv(buf.data(), buf.size(), MPI_INT, 0, 0, MPI_COMM_WORLD,
29             ↪ MPI_STATUS_IGNORE);
30 }
31
32 // While MPI communications are performed here I can do computations as
33 ↪ long as buf is not modified in rank 0, or read/modified in rank 1.
34
35 MPI_Wait(&request, MPI_STATUS_IGNORE);
36 MPI_Finalize();
```

- `MPI_Waitall`, `MPI_Testall` waits or tests completion of all the pending requests.
- `MPI_Waitany`, `MPI_Testany` waits or tests completion of one out of many.
- `MPI_Waitsome`, `MPI_Testsome` waits or tests completion of all the specified requests.
- For arrays of statuses can use `MPI_STATUSES_IGNORE`.
- `MPI_Request_get_status` equivalent to `MPI_Test` but does not free completed requests.

## Syntax

```
1 int MPI_Iprobe(int source, int tag, MPI_Comm comm, int *flag,  
  ↪ MPI_Status *status);  
2  
3 int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status  
  ↪ *status);
```

- Check incoming messages without receiving.
- Immediate variant returns **true** if matching message exists.
- Can be used in combination with a successive `MPI_Get_count` for deducing the size of an incoming message before actually receiving it. Thus, we can allocate a buffer for holding the message.

## Syntax

```
1 int MPI_Barrier(MPI_Comm comm);
```

- Collective communications **must** be called by all processes in the communicator.
- Barrier is hard synchronization.
- Avoid as much as possible.

## Syntax

```
1 int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype, int  
   ↪ sender, MPI_Comm comm);
```

- The **sender** process sends data to every other process.

$P_0$  

$P_1$

$P_2$

$P_3$

## Syntax

```
1 int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype, int  
  ↪ sender, MPI_Comm comm);
```

- The **sender** process sends data to every other process.

$P_0$  

$P_1$  

$P_2$  

$P_3$  

## Syntax

```
1 int MPI_Scatter(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ int sender, MPI_Comm comm);
```

- The **sender** process sends a piece of the data to all processes.
- The sendbuf, sendcount, and sendtype are only relevant on the **sender**.



## Syntax

```
1 int MPI_Scatter(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ int sender, MPI_Comm comm);
```

- The **sender** process sends a piece of the data to all processes.
- The sendbuf, sendcount, and sendtype are only relevant on the **sender**.

 $P_0$   $P_1$   $P_2$   $P_3$  



## Syntax

```
1 int MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ int receiver, MPI_Comm comm);
```

- Every process sends its data to the **receiver** process.
- The `recvbuf`, `recvcount`, and `recvtype` are only relevant on the **receiver**.
- `recvcount` is the size per process, not the total size.

 $P_0$   $P_1$   $P_2$   $P_3$  

## Syntax

```
1 int MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ int receiver, MPI_Comm comm);
```

- Every process sends its data to the **receiver** process.
- The `recvbuf`, `recvcount`, and `recvtype` are only relevant on the **receiver**.
- `recvcount` is the size per process, not the total size.



## Syntax

```
1 int MPI_Allgather(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ MPI_Comm comm);
```

- Every process sends its data to all the other processes.

 $P_0$   $P_1$   $P_2$   $P_3$  

## Syntax

```
1 int MPI_Allgather(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ MPI_Comm comm);
```


- Every process sends its data to all the other processes.



## Syntax

```
1 int MPI_Alltoall(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ MPI_Comm comm);
```


- Every process sends a piece of its data to all the other processes.



## Syntax

```
1 int MPI_Alltoall(const void *sendbuf, int sendcount, MPI_Datatype  
  ↪ sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,  
  ↪ MPI_Comm comm);
```

- Every process sends a piece of its data to all the other processes.



## Syntax

```
1 int MPI_Reduce(const void *sendbuf, void *recvbuf, int count,  
    ↪ MPI_Datatype datatype, MPI_Op op, int receiver, MPI_Comm comm);
```


- Data from all processes are reduced on the **receiver** process.
- Common operations being MPI\_SUM, MPI\_MAX, MPI\_MIN, MPI\_PROD.
- A MPI\_Allreduce variant exists where all processes get the results.
- MPI\_IN\_PLACE can be passed to sendbuf of **receiver**, for *reduce*, and of all processes, for *allreduce*.

 $P_0$   $P_1$   $P_2$   $P_3$  

## Syntax

```
1 int MPI_Reduce(const void *sendbuf, void *recvbuf, int count,  
    ↪ MPI_Datatype datatype, MPI_Op op, int receiver, MPI_Comm comm);
```

- Data from all processes are reduced on the **receiver** process.
- Common operations being MPI\_SUM, MPI\_MAX, MPI\_MIN, MPI\_PROD.
- A MPI\_Allreduce variant exists where all processes get the results.
- MPI\_IN\_PLACE can be passed to sendbuf of **receiver**, for *reduce*, and of all processes, for *allreduce*.





## Syntax

```
1 int MPI_Reduce(const void *sendbuf, void *recvbuf, int count,  
    ↪ MPI_Datatype datatype, MPI_Op op, int receiver, MPI_Comm comm);
```

- Data from all processes are reduced on the **receiver** process.
- Common operations being MPI\_SUM, MPI\_MAX, MPI\_MIN, MPI\_PROD.
- A MPI\_Allreduce variant exists where all processes get the results.
- MPI\_IN\_PLACE can be passed to sendbuf of **receiver**, for *reduce*, and of all processes, for *allreduce*.

 $P_0$   $P_1$  $P_2$  $P_3$

## Syntax

```
1 int MPI_Scan(const void *sendbuf, void *recvbuf, int count,  
  ↪ MPI_Datatype datatype, MPI_Op op, MPI_Comm comm);  
2  
3 int MPI_Exscan(const void *sendbuf, void *recvbuf, int count,  
  ↪ MPI_Datatype datatype, MPI_Op op, MPI_Comm comm);
```

- Performs the prefix reduction on data.
- MPI\_Scan: on process  $i$  contains the reduction of values from processes  $[0, i]$
- MPI\_Exscan: on process  $i$  contains the reduction of values from processes  $[0, i[$ .
- MPI\_IN\_PLACE can be passed to sendbuf.